

Spatial Product Partition Models*

Garrett L. Page[†] and Fernando A. Quintana[‡]

Abstract. When modeling geostatistical or areal data, spatial structure is commonly accommodated via a covariance function for the former and a neighborhood structure for the latter. In both cases the resulting spatial structure is a consequence of implicit spatial grouping in that observations near in space are assumed to behave similarly. It would be desirable to develop spatial methods that explicitly model the partitioning of spatial locations providing more control over resulting spatial structures and be able to better balance local and global spatial dependence. To this end, we extend product partition models to a spatial setting so that the partitioning of locations into spatially dependent clusters is explicitly modeled. We explore the resulting spatial structure and demonstrate its flexibility in accommodating many types of spatial dependencies. We illustrate the method's utility through simulation studies and two applications. Computational techniques with additional simulations are provided in a Supplementary Material file available online.

ppmSuite

Keywords: product partition models, spatial smoothing, spatial clustering, spatial prediction.

1 Introduction

Research dedicated to developing statistical methodologies that in some way incorporate information relating to location has grown rapidly in the last decade. In fact, spatial methods are now available in essentially all areas of statistics and have been developed to accommodate both areal (lattice) and geo-referenced data. The principal motivation in developing these methods is to produce inference and predictions that take into account the spatial dependence that is believed to exist among observations. The end result is typically a smoothed map for areal data or a predictive map for geo-referenced data. These maps are frequently produced by implicitly performing a type of spatial grouping that carries out the intuitively appealing notion that responses measured at locations near in space have similar values. Because the grouping is implicit, the partition of locations is not directly modeled but is a consequence of model choices (e.g., neighborhood structure or covariance function). For areal data this can lead to spatial correlation structures that are counter-intuitive (Wall 2004). Additionally, it is common that the smoothed or predictive maps are global in nature in that methods are not flexible enough to capture local deviations from an overall spatial structure.

*Related articles: DOI: [10.1214/15-BA971A](https://doi.org/10.1214/15-BA971A), DOI: [10.1214/15-BA971B](https://doi.org/10.1214/15-BA971B), DOI: [10.1214/16-BA971C](https://doi.org/10.1214/16-BA971C); rejoinder at DOI: [10.1214/16-BA971REJ](https://doi.org/10.1214/16-BA971REJ).

[†]Departamento de Estadística, Pontificia Universidad Católica de Chile, Santiago, Chile,
page@mat.puc.cl

[‡]Departamento de Estadística, Pontificia Universidad Católica de Chile, Santiago, Chile,
quintana@mat.uc.cl

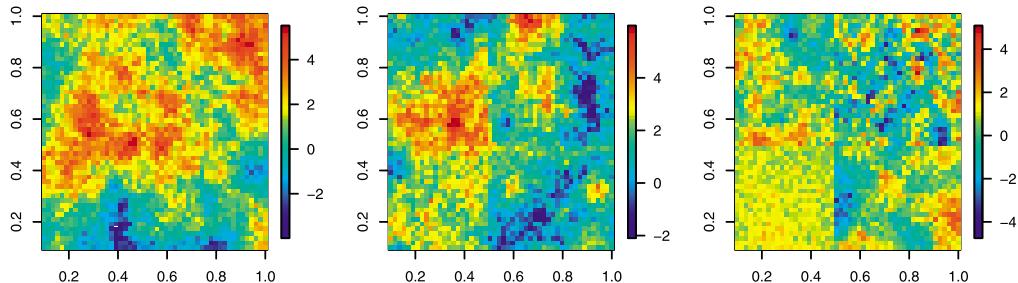


Figure 1: Synthetic spatial fields. From left to right, the graphs display random fields that become progressively more local.

Figure 1 provides a synthetic example of local vs. global spatial dependence. The three plots were generated using a Gaussian process employing an exponential covariance function. From left to right the random fields become increasingly more local. The left plot displays one spatial process over the entire domain that has expectation 0, nugget 0.1, partial sill 2, and effective range 6 (see Banerjee et al. 2015, Chapter 2 for more details). The second plot is generated with the same covariance function, but the field is partitioned into four rectangles with each being assigned a specific constant mean $(1, -0.5, 0.25, -1)$, thus inducing a small amount of local structure. The right plot is the most local of the three as each rectangle is a realization from a unique spatial process that has expectation 0 and a cluster specific partial sill $(1, 2, 3, 4)$ and effective range $(0.5, 10, 5, 20)$. Methods able to flexibly capture these three structures would certainly be appealing and developing these types of methods is the primary focus of this paper.

Our approach is to develop a class of priors based on product partition models (PPM, Hartigan 1990) that directly model the grouping of locations into spatially dependent clusters. Making the PPM location dependent is necessary in a spatial setting because if not, then locations that are very far apart could possibly be assigned to the same cluster with high probability. As a consequence, the marginal correlation between observations far apart could be stronger than that of observations near each other, which runs counter to correlation structures often desired in spatial modeling. As will be seen, PPM's are a very attractive way to partition spatial units as they are extremely flexible in accommodating different types of spatial clusters.

The method we develop is able to adapt to the three scenarios described in Figure 1 by incorporating spatial information in two ways. The first is via a prior on the partitioning of locations using PPM ideas. The second is through the likelihood either directly or hierarchically. If spatial structure is not built in the likelihood, the spatial PPM will marginally induce local spatial dependence among observations. As an aside, apart from more accurately modeling spatial phenomena, considering local spatial dependence potentially provides large computational gains as covariance matrices are considerably smaller.

Spatial methods now have a large presence in the statistical literature. We focus on methods that incorporate spatial dependence flexibly. For a general overview of spatial methods see Gelfand et al. (2010), Banerjee et al. (2015), or Schabenberger and Gotway (2005).

Identifying spatial clustering is common in spatial point processes (Diggle 2014). That said, from a modeling standpoint, the analysis goals are completely different from those we consider. Image segmentation is an extensively studied area that we do not attempt to fully survey here. We do mention the spatial distance dependent Chinese restaurant process of Ghosh et al. (2011) (a spatial extension of the distance dependent Chinese restaurant process of Blei and Frazier 2011) as they develop a process that produces a non-exchangeable distribution on location dependent partitions through a distance dependent decay function. Though there are similarities, our approach is model-based and therefore provides measures of uncertainty regarding inferences and predictions.

Gelfand et al. (2005) developed a spatial Dirichlet process (DP) by taking atoms associated with Sethuraman (1994)'s stick-breaking random measure construction as independent random fields. Duan et al. (2007) generalized the spatial DP through a type of multivariate stick-breaking in which individual sites could possibly arise from unique surfaces introducing a type of local spatial modeling. Both spatial DP processes require replication. Griffin and Steel (2006) developed the ordered dependent DP where stick-breaking weights are randomly permuted according to a latent spatial point process, thus inducing spatial dependence. Petrone et al. (2009) developed a DP that pieces together functions and applied it to a spatial field. Reich and Bondell (2011) use a DP to model locations directly resulting in spatially referenced clusters. All of these methods induce a marginal distribution on partitions through the introduction of latent cluster labels.

Somewhat related to the spatial DP and operationally similar to what we introduce is the spatial stick-breaking process of Reich and Fuentes (2007) and the logistic stick-breaking process of Ren et al. (2011) (both of which are in some sense special cases of kernel stick-breaking process of Dunson and Park 2008). Both stick-breaking processes induce spatial dependence via kernel functions that allow stick-breaking weights to change with space. A related probit stick-breaking prior for spatial dependence was recently proposed in Papageorgiou et al. (2014).

Other authors have employed DP type methods to areal data resulting in a more flexible (local) neighborhood structure (Li et al. 2015, Lee et al. 2014). Kang et al. (2014) created a local conditional autoregressive (CAR) model to accommodate local spatial residual.

Even though all the previously mentioned nonparametric Bayes based methods may have some inferential similarities or are at least operationally similar to what we are proposing, they are fundamentally different. We do not introduce any notion of a random probability measure. Therefore, we are not bound to an induced marginal model on partitions available from the DP (though this particular model is certainly available as a special case). Instead we directly model the spatially dependent partition using a PPM. Doing so provides much more control over the partitioning of spatial units into clusters.

From a disease mapping perspective, Denison and Holmes (2001) consider spatial clustering by first selecting cluster centroids and using tessellation ideas of Lawson and

Denison (2002) to determine cluster memberships. This requires employing Reversible Jump MCMC and produces spatial clusters that are necessarily convex. Knorr-Held and Raßer (2000) cluster areal units via a distance measure that is based on shared boundaries. Hegarty and Barry (2008) employ a PPM to model partitions of areal units, though they do not explore the spatial properties of their model and are restricted to a very specific setting. We aim to propose a very general methodology that is flexible in accommodating many types of spatial dependencies. In fact, we will show that once a model for the partition has been specified, the sky is the limit in terms of how spatial dependence can be incorporated in other parts of the model.

The remainder of the article is organized as follows. In Section 2, we provide some preliminaries on PPM's and a bit of discussion on spatial clustering. Section 3 details spatial extensions of the PPM and investigates spatial properties. Section 4 contains a small simulation study and a Chilean education data application. We make some concluding remarks in Section 5. Lastly, the Supplementary Material file available online contains computational details along with additional simulations and applications (Page and Quintana 2015b).

2 Preliminaries

We provide background to PPM's and a bit of discussion motivating our view of spatial clusters.

2.1 Preliminaries of Product Partition Model

PPMs were first introduced by Hartigan (1990) and have since been extended to include covariates (Müller et al. 2011 and Park and Dunson 2010) and correlated parameters (Monteiro et al. 2011). They have been employed in applications ranging from change point analysis (Barry and Hartigan 1992) to functional clustering (Page and Quintana 2015a) among others. Since PPMs are central to our approach of carrying out spatial clustering, we briefly present them here.

Consider n distinct locations denoted by $\mathbf{s}_1, \dots, \mathbf{s}_n$. The \mathbf{s}_i are quite general in that they can be latitude and longitude values or in the case of areal data they could define a neighborhood structure. Let $\rho_n = \{S_1, \dots, S_{k_n}\}$ denote a partitioning (or clustering) of the n locations into k_n subsets such that $i \in S_h$ implies that location i belongs to cluster h . Alternatively, we will denote cluster membership using c_1, \dots, c_n where $c_i = h$ implies $i \in S_h$. Then the PPM prior for ρ is simply

$$Pr(\rho) \propto \prod_{h=1}^{k_n} C(S_h), \quad (1)$$

where $C(S_h) \geq 0$ for $S_h \subset \{1, \dots, n\}$ is a cohesion function that measures how likely elements of S_h are clustered *a priori*. A popular cohesion function that connects (1) to the marginal prior distribution on partitions induced by a Dirichlet process (DP) is $C(S) = M \times \Gamma(|S|)$. This cohesion produces a PPM that encourages partitions with a small number of large clusters and also a few smaller clusters (the rich get richer

property) and will be helpful in avoiding the creation of many singleton clusters when extending PPMs to a spatial setting. Eventually we will consider a response and covariate vector measured at each location which will be denoted by $y(\mathbf{s}_i)$ and $\mathbf{x}(\mathbf{s}_i)$, respectively. Finally, it will be necessary to make reference to partitioned location and response vectors which we denote by $\mathbf{s}_h^* = \{\mathbf{s}_i : i \in S_h\}$ and $\mathbf{y}_h^* = \{y(\mathbf{s}_i) : i \in S_h\}$.

2.2 Spatial Clustering

Before proceeding, we expound on the term “spatial cluster” and make its definition used in this paper concrete (for more discussion on the subject of spatial clusters see Lawson 2013, Chapter 6). Typically, clustering attempts to group or partition individuals or experimental units based on some measured response variable. Therefore, the resulting partition consists of clusters whose members are fairly homogenous with respect to the measured response. How cluster boundaries are defined (e.g., elliptical, convex) is crucial to the resulting partition and to our knowledge no universally agreed upon definition exists. When in addition to a measured response, the proximity of individuals or experimental units influences the partitioning of individuals, then we refer to these clusters as “spatial”.

If spatial structure exists among the realizations of some response variable measured at various locations, then the values measured at locations near each other should be more similar than those that are far apart. However, this does not exclude the possibility of two individuals far apart producing similar responses. Clustering in the absence of spatial information would group these two individuals together (as would be the case in a non-spatial PPM). From a spatial perspective it seems more natural that locations far apart belong to different clusters. That is, spatial clusters should be in some sense “local” in that locations that belong to the same cluster should share a boundary for areal data (or comply with some other neighborhood structure) or attain a pre-determined minimum distance with other members of the cluster for geo-referenced data. We make this concrete with the following definition.

Definition 1. Consider \mathbf{s}_h^* corresponding to cluster $S_h \subset \{1, \dots, n\}$ and let $d(\cdot, \cdot)$ be a metric in the space of spatial coordinates. We say that cluster S_h is spatially connected if there does not exist $\mathbf{s}_{i'} \notin \mathbf{s}_h^*$ and $\mathbf{s}_i \in \mathbf{s}_h^*$ such that for all $\mathbf{s}_j \in \mathbf{s}_h^*$, $d(\mathbf{s}_{i'}, \mathbf{s}_i) < d(\mathbf{s}_i, \mathbf{s}_j)$. A partition will be called spatially connected if all of its clusters are spatially connected.

Our vision of spatial clusters does not necessarily partition the spatial domain into disjoint sets. Because clusters possibly depend on variables other than location, it is possible that two clusters exist in the same geographical region. The presence of these “stacked” clusters seems common and a perk of the methodology we develop.

3 Methodological Development

We now detail spatial extensions to the basic PPM (here after referred to as sPPM) and investigate cluster membership probabilities. Also, we show that combining sPPM

with likelihoods (that potentially include spatial information) produce marginal spatial structures with appealing properties (e.g., non-stationary) and balance local and global structure. As both cluster membership probabilities and correlations depend on the cohesion function we propose a few reasonable candidates.

3.1 Cohesion Functions

Extending the PPM to incorporate spatial information requires making the cohesion of (1) a function of location. With this in mind, consider

$$Pr(\rho) \propto \prod_{h=1}^{k_n} C(S_h, \mathbf{s}_h^*). \quad (2)$$

Notice that (2) is structurally similar to Park and Dunson 2010's approach to extending the PPM to incorporate covariates. Defining a cohesion function that only admits spatially connected partitions is conceptually straightforward. For example, one could employ

$$C(S, \mathbf{s}_h^*) = \begin{cases} M \times \Gamma(|S|) & \text{if } S \text{ is spatially connected,} \\ 0 & \text{otherwise,} \end{cases}$$

where $M \times \Gamma(|S|)$ is used to favor a small number of large clusters with the number of clusters being regulated by M . A cohesion function defined in this way places zero prior mass on partitions that are not spatially connected. Although this definition is intuitively appealing, it is particularly challenging to implement from a computational stand point and can only realistically be considered for a small number of locations. Therefore, we suggest considering cohesion functions that assign small probabilities to partitions with clusters that are not spatially connected. A nice feature of the sPPM is that there are many ways in which this can be carried out and we introduce four reasonable candidates. Subsequently, we study the spatial properties of each one.

We first approach the construction of (2) from a spatial modeling perspective. That is, distances between the $\mathbf{s}_i \in \mathbf{s}_h^*$ should influence partition probabilities. To this end, we employ tessellation ideas found in Denison and Holmes (2001) in that distances to a cluster centroid are used to penalize partitions with clusters that are spatially disperse. Now, let $\bar{\mathbf{s}}_h$ denote the centroid of cluster S_h whose coordinates are calculated using $\bar{s}_{hk} = 1/n_h \sum_{i \in S_h} s_{ik}$ for $k = 1, 2$ and $n_h = |S_h|$, and let $\mathcal{D}_h = \sum_{i \in S_h} d(\mathbf{s}_i, \bar{\mathbf{s}}_h)$ denote the sum of all centroid distances (unless otherwise stated we use the Euclidean norm $\|\cdot\|$). Defining the cohesion as a decreasing function of \mathcal{D}_h would certainly produce small local clusters. Unfortunately, cohesions that favor clusters with small \mathcal{D}_h would also produce partitions with many singletons. To prevent this, we make the cohesion a function of $M \times \Gamma(|S_h|)$ in addition to \mathcal{D}_h . This requires using $\Gamma(\mathcal{D}_h)$ to ensure that $|S_h|$ and \mathcal{D}_h are weighted similarly. Now, since \mathcal{D}_h is continuous on \mathbb{R}^+ and $\Gamma(\cdot)$ is not monotone increasing on $[0, 1]$ we consider $\Gamma(\mathcal{D}_h)\mathbb{I}[\mathcal{D}_h \geq 1] + \mathcal{D}_h\mathbb{I}[\mathcal{D}_h < 1]$. Finally, to provide a bit more control over the penalization of distances, we introduce a user

supplied tuning parameter (α) resulting in the following cohesion function

$$C_1(S_h, \mathbf{s}_h^*) = \begin{cases} \frac{M \times \Gamma(|S_h|)}{\Gamma(\alpha \mathcal{D}_h) \mathbb{I}[\mathcal{D}_h \geq 1] + (\mathcal{D}_h) \mathbb{I}[\mathcal{D}_h < 1]} & \text{if } |S_h| > 1, \\ M & \text{if } |S_h| = 1. \end{cases} \quad (3)$$

We set $C_1(S_h, \mathbf{s}_h^*) = M$ for $|S_h| = 1$ to avoid issues associated with $\mathcal{D}_h = 0$. Notice that since all $\mathbf{s}_1, \dots, \mathbf{s}_n$ are distinct $\mathcal{D}_h = 0 \iff |S_h| = 1$. Further, when $|S_h| = 1$, $M \times \Gamma(|S_h|) = M$ justifying in a sense setting the cohesion to M when $|S_h| = 1$.

The second cohesion function we consider provides a hard cluster boundary and for some pre-specified $a > 0$ has the following form

$$C_2(S_h, \mathbf{s}_h^*) = M \times \Gamma(|S_h|) \times \prod_{i,j \in S_h} \mathbb{I}[\|\mathbf{s}_i - \mathbf{s}_j\| \leq a]. \quad (4)$$

Once again, $M \times \Gamma(|S_h|)$ is included to control the number of singleton clusters. This cohesion is amenable to neighborhood structures commonly employed in areal data modeling. Instead of $\mathbb{I}[d(\mathbf{s}_i, \mathbf{s}_j) \leq a]$, one could use $\mathbb{I}[i \sim j]$ where $i \sim j$ indicates that \mathbf{s}_i and \mathbf{s}_j are neighbors according to some neighborhood structure. If a data dependent neighborhood structure is desired, one could introduce auxiliary variables in the cohesion and employ ideas similar to those found in Kang et al. (2014).

An sPPM under C_1 and C_2 produces a completely valid joint distribution over partitions that is quite general. In fact, since the cohesions are functions of not only $|S_h|$ but also of \mathbf{s}_h^* , sPPM relaxes exchangeability assumptions. However, for this same reason sPPM under C_1 and C_2 does not inherit the PPM's (1) property of being coherent across sample sizes. That is, $Pr(\rho_n) \neq \sum_{h=1}^{k_n+1} Pr(\rho_n, c_{n+1} = h)$. This is easily seen as the location of s_{n+1} influences $Pr(\rho_n, c_{n+1} = j)$. Although this does not change the fact that the sPPM produces a valid joint distribution over partitions, for computational purposes it is sometimes desirable to have coherence across sample sizes. To retain this property one would need to "marginalize" over all possible locations. This was considered in detail in Müller et al. (2011) (and also mentioned in Park and Dunson 2010) when making a PPM covariate dependent. We employ ideas developed in Müller et al. (2011) in a spatial setting by essentially treating the \mathbf{s}_i as covariates and using the following cohesion

$$C_3(S_h, \mathbf{s}_h^*) = M \times \Gamma(|S_h|) \times \int \prod_{i \in S_h} q(\mathbf{s}_i | \boldsymbol{\xi}_h) q(\boldsymbol{\xi}_h) d\boldsymbol{\xi}_h. \quad (5)$$

Müller et al. (2011) call $q(\mathbf{s}_i | \boldsymbol{\xi}_h) q(\boldsymbol{\xi}_h)$ an auxiliary model and even though in theory $q(\cdot | \cdot)$ and $q(\cdot)$ could represent any combination of functions, for reasons detailed in Müller et al. (2011) they will denote a conjugate pair of densities. Therefore, C_3 favors partitioned location vectors (\mathbf{s}^*) that produce large marginal likelihood values. We emphasize, however, that we are not assuming the \mathbf{s}_i 's to be random, we are simply employing the conjugate model as a means to measure spatial proximity and encourage co-clustering of locations that are near each other. Both areal and point referenced data can be considered when C_3 is employed though we focus on the later where a con-

jugate Gaussian/Gaussian–Inverse-Wishart model would be appropriate. In this case $\xi = (\mathbf{m}, \mathbf{V})$ would denote a mean and covariance, $q(\mathbf{s}|\xi) = N(\mathbf{s}|\mathbf{m}, \mathbf{V})$ a bivariate Gaussian density, and $q(\xi) = NIW(\mathbf{m}, \mathbf{V}|\boldsymbol{\mu}_0, \kappa_0, \nu_0, \boldsymbol{\Lambda}_0)$ a bivariate Normal–Inverse-Wishart density. In what follows, we will occasionally refer to C_3 as the auxiliary cohesion. Finally, as in the previous two cohesions, $M \times \Gamma(|S_h|)$ is included so that partitions containing many singletons are discouraged *a priori*.

The fourth and final cohesion that we consider is similar to what Quintana et al. (In press) call the “double dipper”. It has the same form as C_3 , but instead of employing a prior predictive conjugate model, a posterior predictive conjugate model is used. Therefore, C_4 has the following form

$$C_4(S_h, \mathbf{s}_h^*) = M \times \Gamma(|S_h|) \times \int \prod_{i \in S_h} q(\mathbf{s}_i|\xi_h) q(\xi_h|\mathbf{s}_h^*) d\xi_h. \quad (6)$$

Since the posterior predictive is typically more peaked than the prior predictive, C_4 puts more weight on partitions that are local. Once again both areal and point-referenced data are possible, but in what follows we focus on point-referenced and use $N_2(\mathbf{s}_i|\mathbf{m}_h, \mathbf{V}_h) NIW(\mathbf{m}_h, \mathbf{V}_h|\mathbf{s}_h^*)$ as the conjugate model.

Before proceeding we briefly discuss M ’s role in sPPM. It is fairly well known that the expected number of clusters *a priori* under the DP induced partition probability distribution is approximately $M \log(n)$. Thus the number of clusters grows slowly as n increases which favors partitions with a small number of large clusters (motivating its inclusion in $C_1 - C_4$). However, when $M \times \Gamma(|S_h|)$ is coupled with distance penalties, it is not clear how the number of expected clusters *a priori* grows as a function of M (or n). We explore this using a small simulation study in the next section.

3.2 Cluster Assignment Probabilities

To investigate how distance influences partition (cluster membership) probabilities we consider the very simple case of $n = 2$. In this context only two possible partitions exist: $(\{1, 2\})$ and $(\{1\}, \{2\})$. Table 1 provides $Pr(\rho = \{1, 2\})$ for each of the cohesion functions along with the limiting probabilities as $d(\mathbf{s}_1, \mathbf{s}_2) \rightarrow 0$ and $d(\mathbf{s}_1, \mathbf{s}_2) \rightarrow \infty$. To simplify calculations, for the auxiliary and double dipping similarity functions we use $\boldsymbol{\mu}_0 = \bar{\mathbf{s}}_h$, $\kappa_0 = 1$, $\nu_0 = 2$, and $\boldsymbol{\Lambda}_0 = \text{diag}(1, 1)$. Also, in Table 1, $\mathbf{S} = \sum_{i \in S_h} (\mathbf{s}_i - \bar{\mathbf{s}}_h)(\mathbf{s}_i - \bar{\mathbf{s}}_h)'$.

From Table 1 it can be seen that for all four cohesions the probability that both locations are members of the same cluster approaches zero as distance between the two locations increases (a quality that is desirable). However, only C_1 displays the desirable property (from a spatial perspective) that the probability of two locations being assigned to the same cluster approaches one as the distance between the two locations decreases. This limiting probability for the other three cohesion functions depends on M and other tuning parameter choices. A slightly more sophisticated example that further explores partition probabilities is provided in the Supplementary Material (Page and Quintana 2015b).

To further explore probabilities of locations co-clustering we drew 5000 samples from the sPPM on a 10×10 regular grid and calculated empirical pairwise probabilities.

Cohesion	$Pr(\{1, 2\})$	$d(s_1, s_2) \rightarrow 0$	$d(s_1, s_2) \rightarrow \infty$
		$Pr(\{1, 2\})$	$Pr(\{1, 2\})$
$C_1(S_h, s_h^*)$	$\frac{1}{1 + M\{\Gamma(\alpha D_h)I[D_h \geq 1] + D_h I[D_h < 1]\}}$	1	0
$C_2(S_h, s_h^*)$	$\frac{I[d(s_1, s_2) \leq a]}{I[d(s_1, s_2) \leq a] + M}$	$\frac{1}{1 + M}$	0
$C_3(S_h, s_h^*)$	$\frac{1}{1 + 2M \Lambda_0 + S ^{3/2}}$	$\frac{1}{1 + 2M}$	0
$C_4(S_h, s_h^*)$	$\frac{81 \Lambda_0 + S ^2}{81 \Lambda_0 + S ^2 + 10M \Lambda_0 + 2S ^3}$	$\frac{81}{81 + 10M}$	0

Table 1: Prior Partition Probabilities for ($\{1, 2\}\right).$

Figure 2 displays the results. We briefly note that since sPPM under C_1 and C_2 is not coherent across sample sizes, care must be taken when generating samples from the prior and we use self-normalized importance sampling (Robert and Casella 2010, chap 3). In Figure 2, M is set to 0.1 for C_1 and C_2 and $M = 1$ for C_3 and C_4 . For C_2 we set $a = 1.77$ which is the median distance among all pairwise distances, and the tuning parameters associated with C_1 , C_3 and C_4 are those used previously. From Figure 2 it appears that C_1 and C_4 are similar in how distance penalizes cluster membership. C_3 allows locations fairly far apart to have positive probability of being members of the same cluster. The cut-off boundary for cluster membership associated with C_2 is clearly shown.

To better understand the influence that the scale parameter (M) has on partition probabilities available from sPPM we provide results from a small simulation study. In the the study 5000 partitions were drawn from the sPPM prior for each of the four cohesion. The spatial configurations used are regular 10×10 , 15×15 and 20×20 grids resulting in 100, 225, and 400 spatial locations. (We also considered the spatial configuration found in the application of Section 4.2 but results were similar and so are not provided.) The tuning parameters are set to the same values as used previously except that both $\alpha = 1, 2$ are considered for C_1 . The results are provided in Table 2. Each column of the table represents an average over the 5000 prior draws. Values under the column $E(k_n)$ are the expected number of clusters in ρ , #sing denotes the number of singletons clusters and $\max |S_j|$ denotes the membership size of the largest cluster. Notice that setting $a = 1.77$ for C_2 forces the sPPM to have at least 10 clusters. Also, as expected, setting $\alpha = 2$ results in C_1 producing more clusters. The number of clusters associated with C_1 , C_2 , and C_4 grow at a faster rate than $M \log(n)$ while C_3 grows at a slower rate. The number of singleton clusters is also very reasonable for $M \leq 1$.

3.3 Modeling Spatial Structure via the Likelihood and Prior

Given ρ , the sky is the limit on how spatial dependence might be modeled via the likelihood. A completely valid modeling strategy would be to assume independent ob-

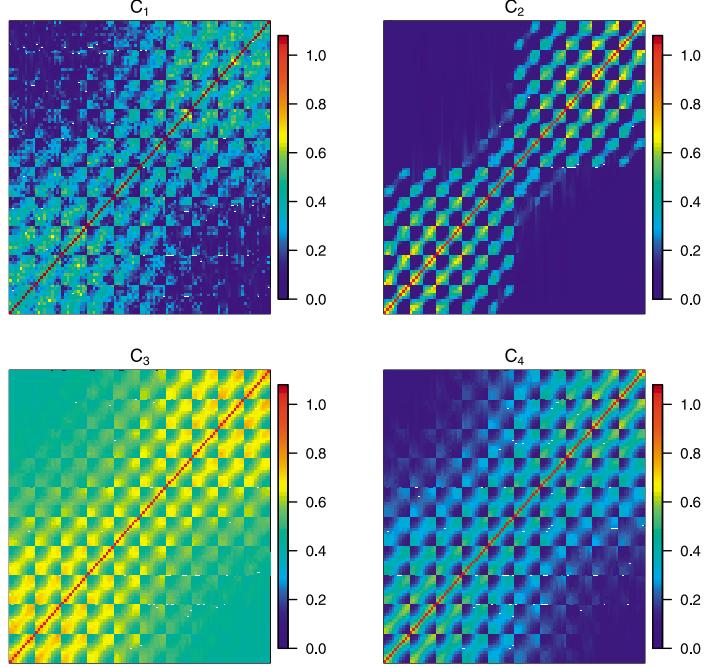


Figure 2: Pairwise probability matrix of two locations belonging to the same cluster for a 10×10 regular grid. $M = 0.1$ for each cohesion.

servations given ρ . In this case, all spatial dependence would originate from the spatial clustering produced by the sPPM. Alternatively, global or cluster specific spatial structure may be included in the likelihood producing much richer spatial structures.

To explore spatial dependence available from the sPPM, we consider correlations among two observations as distance between them either increases to ∞ or decreases to 0. This is done under a few likelihood models for each of the cohesions. In the absence of spatial dependence in the likelihood, the basic model is

$$\begin{aligned} f(\mathbf{y}|\rho) &= \prod_{h=1}^{k_n} f_h(\mathbf{y}_h^*), \\ Pr(\rho) &\propto \prod_{h=1}^{k_n} C(S_h, \mathbf{s}_h^*), \end{aligned} \quad (7)$$

where $\mathbf{y} = (y(\mathbf{s}_1), \dots, y(\mathbf{s}_n))$ and $f_h(\mathbf{y}_h^*) = \int \prod_{i \in S_h} f(y(\mathbf{s}_i)|\boldsymbol{\theta}) dG_0(\boldsymbol{\theta})$ with $f(\cdot|\boldsymbol{\theta})$ denoting the likelihood and G_0 a prior on $\boldsymbol{\theta}$. Alternatively, the model can be written hierarchically using cluster labels c_1, \dots, c_n in the following way

$$\begin{aligned} y(\mathbf{s}_i) | \boldsymbol{\theta}, c_i &\stackrel{ind}{\sim} f(\theta_{c_i}^*), \text{ for } i = 1, \dots, n, \\ \theta_\ell^* &\stackrel{iid}{\sim} G_0, \text{ for } \ell = 1, \dots, k_n, \end{aligned} \quad (8)$$

M	Method	n = 100			n = 225			n = 400		
		$E(k_n)$	#sing	$\max S_j $	$E(k_n)$	#sing	$\max S_j $	$E(k_n)$	#sing	$\max S_j $
10^{-5}	$C_{1_{\alpha=1}}$	1.00	0.00	100.00	1.00	0.00	224.99	1.01	0.00	399.99
	$C_{1_{\alpha=2}}$	3.91	0.03	37.06	4.61	0.01	66.85	4.98	0.00	106.92
	C_2	10.08	0.82	18.18	11.63	0.68	39.11	13.06	0.64	67.59
	C_3	1.00	0.00	100.00	1.00	0.00	225.00	1.00	0.00	400.00
	C_4	1.00	0.00	99.98	1.00	0.00	224.99	1.00	0.00	399.96
10^{-4}	$C_{1_{\alpha=1}}$	1.01	0.01	99.96	1.03	0.02	224.93	3.00	0.00	345.00
	$C_{1_{\alpha=2}}$	4.58	0.04	31.04	5.40	0.00	57.28	7.00	0.00	80.02
	C_2	10.11	0.81	18.20	11.65	0.68	39.13	13.08	0.64	67.53
	C_3	1.00	0.00	99.99	1.00	0.00	224.98	1.00	0.00	399.92
	C_4	1.00	0.00	99.97	1.00	0.00	224.90	1.00	0.00	399.86
10^{-3}	$C_{1_{\alpha=1}}$	1.16	0.03	99.37	2.17	0.00	141.19	2.77	0.00	227.96
	$C_{1_{\alpha=2}}$	5.50	0.00	25.76	6.76	0.00	49.19	8.08	0.00	68.10
	C_2	10.10	0.82	18.15	11.65	0.68	39.15	13.05	0.64	67.49
	C_3	1.00	0.00	99.93	1.00	0.00	224.85	1.01	0.00	399.52
	C_4	1.02	0.00	99.62	1.02	0.00	224.00	1.02	0.00	398.27
10^{-2}	$C_{1_{\alpha=1}}$	3.00	0.01	55.99	3.18	0.00	95.76	3.00	0.00	151.00
	$C_{1_{\alpha=2}}$	8.43	0.03	20.62	9.51	0.02	39.33	12.93	0.00	53.83
	C_2	10.17	0.84	18.13	11.72	0.70	39.10	13.20	0.65	67.30
	C_3	1.04	0.01	99.22	1.05	0.01	223.42	1.05	0.01	396.73
	C_4	1.16	0.01	96.33	1.17	0.01	217.12	1.19	0.01	385.04
10^{-1}	$C_{1_{\alpha=1}}$	5.91	0.22	30.66	8.87	0.00	46.20	8.50	0.02	83.57
	$C_{1_{\alpha=2}}$	14.12	0.73	13.78	18.98	0.63	22.30	25.03	0.31	32.20
	C_2	10.89	1.00	17.77	12.69	0.89	38.15	14.34	0.85	65.57
	C_3	1.42	0.07	92.84	1.46	0.07	209.11	1.51	0.07	370.28
	C_4	2.22	0.10	76.89	2.40	0.09	171.75	2.52	0.10	304.25
10^0	$C_{1_{\alpha=1}}$	14.96	1.11	14.24	21.66	0.63	22.48	31.03	1.21	31.20
	$C_{1_{\alpha=2}}$	26.50	2.57	7.85	43.98	3.27	11.70	54.80	1.74	16.78
	C_2	17.84	3.19	14.54	22.31	2.96	30.38	26.23	3.00	51.37
	C_3	4.27	0.72	62.99	4.64	0.70	141.88	5.01	0.71	249.55
	C_4	7.70	0.97	35.91	9.17	0.94	76.42	10.22	0.96	132.32
10^1	$C_{1_{\alpha=1}}$	36.51	9.28	7.06	60.10	9.61	10.12	85.77	10.31	13.30
	$C_{1_{\alpha=2}}$	52.34	19.55	4.46	92.38	19.91	6.68	137.61	19.82	7.87
	C_2	46.78	21.86	7.21	70.16	23.34	13.27	92.31	24.77	20.80
	C_3	18.83	6.59	25.10	23.02	6.80	56.47	25.86	6.89	99.96
	C_4	27.72	8.93	12.88	37.99	9.19	25.10	46.30	9.33	41.22

Table 2: Results from prior simulation study. 5,000 partitions were drawn from the sPPM for each of the four cohesions. $E(k_n)$ denotes the expected number of clusters *a priori*, #sing the number of singleton clusters, and $\max|S_j|$ the size of the largest cluster in the partition.

with $\theta_1^*, \dots, \theta_{k_n}^*$ denoting cluster specific parameters so that $\theta_i = \theta_{c_i}^*$. In the spatial setting c_1, \dots, c_n are dependent multinomial latent variables with component probabilities derived from the sPPM.

Spatial structure will be included in the likelihood hierarchically by way of spatial random effects and models (7) and (8) will need to be adjusted accordingly. The spatial random effects can be cluster specific or global. If covariates are available, their relationship to the response can also be modeled as being cluster specific (local) or not (global). To simplify calculations, in what follows we consider a Gaussian likelihood by setting $f(\cdot|\boldsymbol{\theta}) = N(\cdot|\mu, \sigma^2)$.

Correlations Under Local Regression

Proposition 3.1 furnishes the correlation between two observations available from a model that incorporates spatial information in the prior only. Therefore, all spatial structure is completely produced by the sPPM. The proof of this proposition and all subsequent propositions are provided in Appendix A.

Proposition 3.1. *Let $\mathbf{x}(\mathbf{s}_i) = \mathbf{x}_i$ and $y(\mathbf{s}_i) = y_i$ denote a p -dimensional covariate vector and response at location \mathbf{s}_i . Further, let $\beta_1^*, \dots, \beta_{k_n}^*$ denote cluster specific parameters such that $\beta_h^* \stackrel{iid}{\sim} N(\boldsymbol{\mu}, \mathbf{T})$ and assume that ρ and $\{\beta_h^*\}_{h=1}^{k_n}$ are mutually independent. Then under likelihood*

$$y_i | \mathbf{x}_i, c_i, \boldsymbol{\beta}^*, \sigma^2 \sim N(\mathbf{x}'_i \boldsymbol{\beta}_{c_i}^*, \sigma^2) \quad (9)$$

and an sPPM prior for ρ , the marginal correlation between two observations is

$$\text{corr}(y_i, y_j) = \frac{\mathbf{x}'_i \mathbf{T} \mathbf{x}_j}{\sqrt{\mathbf{x}'_i \mathbf{T} \mathbf{x}_i + \sigma^2} \sqrt{\mathbf{x}'_j \mathbf{T} \mathbf{x}_j + \sigma^2}} \Pr(c_i = c_j). \quad (10)$$

When $\mathbf{x}(\mathbf{s}_i) = 1$ for all i (i.e., no covariates are available) and $\beta_h^* \stackrel{iid}{\sim} N(\mu, \tau^2)$, then (10) simplifies to

$$\text{corr}(y_i, y_j) = \frac{\tau^2}{\tau^2 + \sigma^2} \Pr(c_i = c_j). \quad (11)$$

Remark 3.1. Recall that as $d(\mathbf{s}_i, \mathbf{s}_j) \rightarrow \infty$, $\Pr(c_i = c_j) \rightarrow 0$ and therefore $\text{corr}(y_i, y_j) \rightarrow 0$. However, $\text{corr}(y_i, y_j) \not\rightarrow 1$ as $d(\mathbf{s}_i, \mathbf{s}_j) \rightarrow 0$. Although this result does not agree with many spatial covariance functions, it does agree with models that include a nugget effect. Additionally, from a clustering perspective it makes sense that locations allocated to the same cluster are assigned the same latent parameter value, but not the same response value.

To visualize (11) as a function of distance, $d(\mathbf{s}_1, \mathbf{s}_2) = \|\mathbf{s}_1 - \mathbf{s}_2\|$, consider again the case of two locations. In Figure 3, we present correlations that are calculated by fixing $\mathbf{s}_1 = (0, 0)$ and moving \mathbf{s}_2 around in space. We set $\sigma^2 = 0.1$ and $\tau^2 = 1$ which produces $1/1.1 \approx 0.9$ as the maximum correlation. For each cohesion we set $M = 1$ and use the same values for the tuning parameters that were used in Section 3.2. The hard boundary of C_2 is evident as correlations produced by C_2 are either zero or $0.5(1/1.1) \approx 0.45$. The correlations associated with the other three cohesions decrease more smoothly as

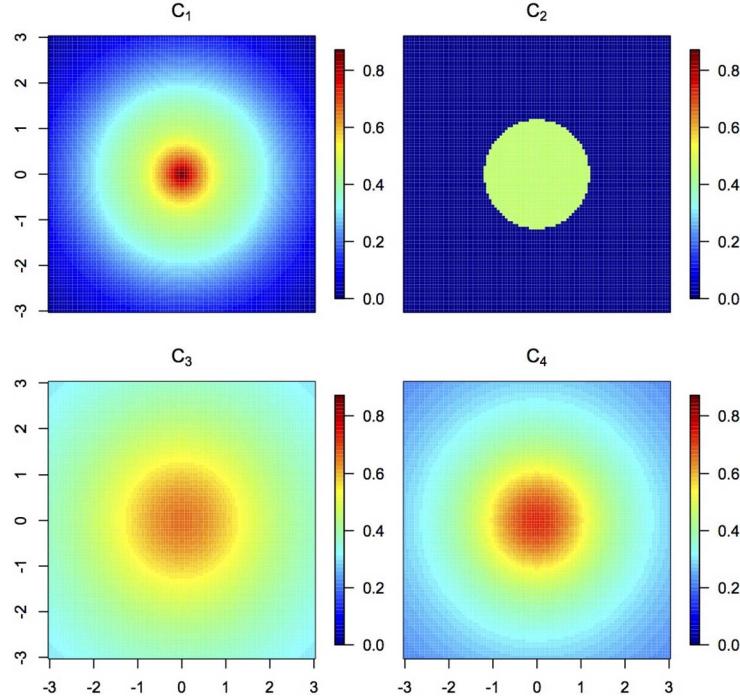


Figure 3: Correlations produced using (11) when two locations are considered. s_1 is set to $(0, 0)$ and s_2 varies. The maximum correlation available is $\tau^2/(\tau^2 + \sigma^2) \approx 0.91$ with $\tau^2 = 1.0$ and $\sigma^2 = 0.1$.

distances between s_1 and s_2 increase. It appears that correlations associated with C_1 decay quicker as distance increases relative to C_3 and C_4 . The correlations associated with C_3 seem to be the most global in the sense that they decay slowly as a function of distance.

In order to consider simultaneous movement between two observations, in Figure 4 $s_1, s_2 \in \mathbb{R}$ (rather than $s_1, s_2 \in \mathbb{R}^2$). Thus what is seen in Figure 4 are correlations associated with $d(s_1, s_2) = |s_1 - s_2|$. Once again the maximum correlation is 1/1.1. Just as in the previous figure, C_2 's hard boundary is evident and C_1 displays the most extreme correlation values. However, perhaps more interesting is the fact that the spatial structures produced by C_3 and C_4 appear to be nonstationary and anisotropic as they are not constant in distance or direction.

Correlations Under Local Regression and Global Spatial Structure

Proposition 3.2. Let x_i , y_i , and $\beta_1^*, \dots, \beta_{k_n}^*$ be as described in Proposition 3.1. Further, Let $\theta = [\theta(s_1), \dots, \theta(s_n)] \sim GP(0, \lambda^2 H(\phi))$ denote an n -dimensional vector of a spatial process where $GP(0, \lambda^2 H(\phi))$ denotes a Gaussian process with covariance func-

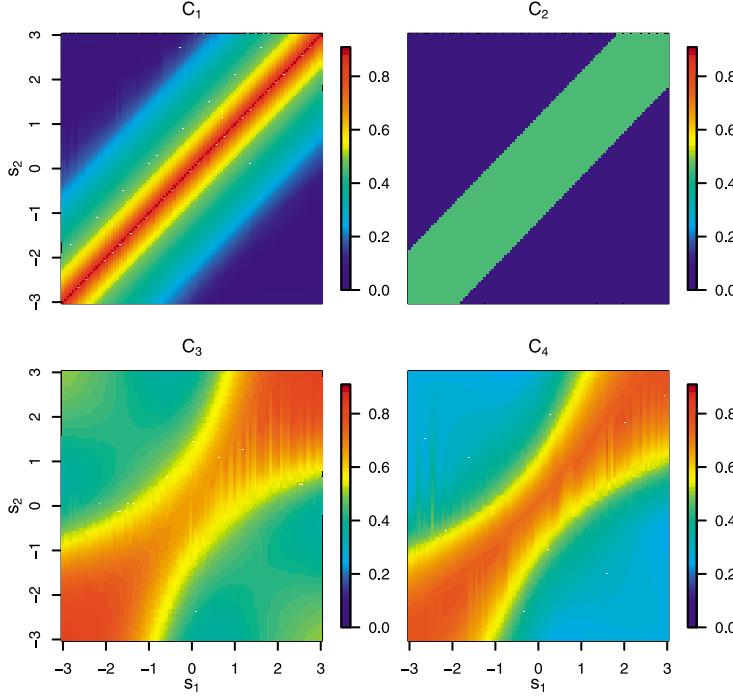


Figure 4: Pairwise correlations calculated using (11) and distances $|s_1 - s_2|$. The maximum correlation available is $\tau^2 / (\tau^2 + \sigma^2) \approx 0.91$ with $\tau^2 = 1.0$ and $\sigma^2 = 0.1$.

tion $H(\phi) : \mathbb{R}^2 \times \mathbb{R}^2 \rightarrow \mathbb{R}$ parametrized by ϕ . Assume that ρ , $\{\beta_h^*\}_{h=1}^{k_n}$, and θ are mutually independent. Then for likelihood

$$y_i | \mathbf{x}_i, \theta_i, \beta^*, c_i, \sigma^2 \sim N(\mathbf{x}'_i \beta_{c_i}^* + \theta_i, \sigma^2) \quad (12)$$

and *sPPM prior on ρ* , the marginal correlation between two observations is

$$\text{corr}(y_i, y_j) = \frac{\lambda^2 (H(\phi))_{i,j} + \mathbf{x}'_j T \mathbf{x}_i \Pr(c_i = c_j)}{\sqrt{\mathbf{x}'_i T \mathbf{x}_i + \lambda^2 + \sigma^2} \sqrt{\mathbf{x}'_j T \mathbf{x}_j + \lambda^2 + \sigma^2}}. \quad (13)$$

When $\mathbf{x}(s_i) = 1$ for all i (i.e., no covariates are available) and $\beta_h^* \stackrel{iid}{\sim} N(\mu, \tau^2)$, then (13) simplifies to

$$\text{corr}(y_i, y_j) = \frac{\lambda^2}{\tau^2 + \lambda^2 + \sigma^2} (H(\phi))_{i,j} + \frac{\tau^2}{\tau^2 + \lambda^2 + \sigma^2} \Pr(c_i = c_j). \quad (14)$$

Correlations are now a function of covariances from the GP and from spatial clustering. Notice that if the variability among cluster means (τ^2) is large relative to σ^2 and λ^2 , then cluster probabilities will be extremely influential in marginal correla-

tions. Consider once again the simple case of two spatial locations. In this scenario, if $d(\mathbf{s}_1, \mathbf{s}_2) \rightarrow \infty$, then $\text{corr}(y_1, y_2) \rightarrow 0$. While as $d(\mathbf{s}_1, \mathbf{s}_2) \rightarrow 0$, then $\text{corr}(y_1, y_2) \rightarrow (\lambda^2 + \tau^2 \Pr(c_1 = c_2)) / (\lambda^2 + \tau^2 + \sigma^2)$. Thus modeling spatial partitions with the sPPM results in decreased correlation for locations that have small probability of being co-clustered and an increase for those that have high probability relative to GP type spatial structures.

Covariances Under Global Regression and Local Spatial Structure

Proposition 3.3. Let \mathbf{x}_i, y_i be as described in Proposition 3.1. Further let $\boldsymbol{\beta} \sim N(\boldsymbol{\mu}, \mathbf{T})$ denote the global regression coefficient and $\boldsymbol{\theta}_h = \{\theta_i : i \in S_h\}$ such that $\boldsymbol{\theta}_h | \lambda_h^{2*}, \phi_h^* \sim GP(0, \lambda_h^{2*} H(\phi_h^*))$. With out loss of generality, order $\boldsymbol{\theta} = (\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_{k_n})$ such that

$$\begin{pmatrix} \boldsymbol{\theta}_1 \\ \vdots \\ \boldsymbol{\theta}_{k_n} \end{pmatrix} \sim N_n \left(\mathbf{0}, \begin{bmatrix} \lambda_1^{2*} H(\phi_1^*) & \dots & \mathbf{0} \\ \vdots & \ddots & \vdots \\ \mathbf{0} & \dots & \lambda_{k_n}^{2*} H(\phi_{k_n}^*) \end{bmatrix} \right). \quad (15)$$

If spatial random effects (15) are combined with likelihood $y_i | \mathbf{x}_i, \theta_i, \boldsymbol{\beta}, c_i, \sigma^2 \sim N(\mathbf{x}'_i \boldsymbol{\beta} + \theta_i, \sigma^2)$ and sPPM is employed to model ρ with ρ , $\boldsymbol{\beta}$, and $\boldsymbol{\theta}$ being mutually independent, then the marginal correlation between two observations is

$$\text{corr}(y_i, y_j) = \frac{\mathbf{x}'_j \mathbf{T} \mathbf{x}_i + \text{cov}^*(\theta_i, \theta_j)}{\sqrt{\sigma^2 + \mathbf{x}'_i \mathbf{T} \mathbf{x}_i + \text{var}^*(\theta_i)} \sqrt{\sigma^2 + \mathbf{x}'_j \mathbf{T} \mathbf{x}_j + \text{var}^*(\theta_j)}}, \quad (16)$$

where

$$\text{cov}^*(\theta_i, \theta_j) = \sum_{h=1}^{k_n} \lambda_h^{2*} (H(\phi_h^*))_{i,j} \Pr(c_i = c_j = h) \quad \text{and} \quad \text{var}^*(\theta_i) = \sum_{h=1}^{k_n} \lambda_h^{2*} \Pr(c_i = h).$$

When $\mathbf{x}(\mathbf{s}_i) = 1$ for all i (i.e., no covariates are available) and $\boldsymbol{\beta} \sim N(\boldsymbol{\mu}, \tau^2)$, then (16) simplifies to

$$\text{corr}(y_i, y_j) = \frac{\tau^2 + \text{cov}^*(\theta_i, \theta_j)}{\sqrt{\sigma^2 + \tau^2 + \text{var}^*(\theta_i)} \sqrt{\sigma^2 + \tau^2 + \text{var}^*(\theta_j)}}. \quad (17)$$

It is interesting to note that covariances are weighted averages of all cluster-specific covariances with weights depending on distance. This type of spatial correlation structure is clearly nonstationary and nonisotropic.

4 Simulation Study and Examples

Except for very specific examples, the discussion to this point has been fairly generic with the idea of explaining different modeling approaches under a general framework. Now we provide more concrete illustrations by way of a small simulation study, and

applications to Chilean education data and to a well-known dataset on scallop catch in the New York/New Jersey Bight. Additional simulations are provided in the Supplementary Material (Page and Quintana 2015b). The simulation studies and applications will require making some specific modeling assumptions but still within the general class of models thus far presented. To make methods invariant to scale of location, in the simulations and applications that follow we standardize $\mathbf{s}_1, \dots, \mathbf{s}_n$ to have mean zero and unit variance. Fitting the models that will be described is a straightforward MCMC exercise. The algorithm we employ is based on Neal (2000)'s algorithm number 8 and details are provided in the Supplementary Material (Page and Quintana 2015b).

4.1 Simulation Study

We conduct a small simulation study to explore sPPM's ability to recover partitions, make predictions and assess its goodness-of-fit performance. This is done by specifying the following model:

$$\begin{aligned} y(\mathbf{s}_i) | x(\mathbf{s}_i), c_i, \mu_{c_i}^*(\mathbf{s}_i), \sigma^2 &\stackrel{iid}{\sim} N(\mu_{c_i}^*(\mathbf{s}_i) + x(\mathbf{s}_i)\beta, \sigma^2), \\ \sigma &\sim UN(0, 10), \beta \sim N(0, 10^2), \\ \mu_h^*(\mathbf{s}_i) &\stackrel{iid}{\sim} N(\mu_0, \sigma_0^2) \text{ for } h = 1, \dots, k_n, \text{ and} \\ \mu_0 &\sim N(0, 10^2), \sigma_0 \sim UN(0, 10), \\ \{c_i\}_{i=1}^n &\sim sPPM. \end{aligned} \tag{18}$$

Here after this procedure will be referred to as the Conditional Model with Prior Spatial Structure (CPS). For C_1 we consider $\alpha = 1$ and $\alpha = 2$ and use the same tuning parameter values as in Section 3.2 for the other three cohesion functions. To the CPS we compare the spatial stick breaking (SSB) process found in Reich and Fuentes (2007) and a common spatial regression model (SR). More precisely,

1. The SR model refers to $y(\mathbf{s}_i) | \mathbf{x}(\mathbf{s}_i), \boldsymbol{\beta}, \theta(\mathbf{s}_i) \sim N(\mathbf{x}'(\mathbf{s}_i)\boldsymbol{\beta} + \theta(\mathbf{s}_i), \sigma^2)$ with $\mathbf{x}'(\mathbf{s}_i) = (1, x(\mathbf{s}_i))$, $\boldsymbol{\beta} = (\beta_0, \beta_1) \sim N_2(\mathbf{0}, 10^2 \mathbf{I})$, $[\theta(\mathbf{s}_1), \dots, \theta(\mathbf{s}_n)] \sim GP(0, \lambda^2 H(\phi))$, and $\sigma^2 \sim IG(a, b)$.
2. Given cluster labels $\{c_i\}_{i=1}^n$, SSB can be expressed as $y(\mathbf{s}_i) | x(\mathbf{s}_i), c_i, \mu_{c_i}^*(\mathbf{s}_i), \sigma^2 \sim N(\mu_{c_i}^*(\mathbf{s}_i) + x(\mathbf{s}_i)\beta, \sigma^2)$ where $c_i \sim \text{Categorical}(p_1(\mathbf{s}_i), \dots, p_m(\mathbf{s}_i))$ with $p_j(\mathbf{s}) = w_j(\mathbf{s}) V_j \prod_{k < j} [1 - w_k(\mathbf{s}) V_k]$ for $V_j \stackrel{iid}{\sim} \text{Beta}(1, M)$. The $w_j(\mathbf{s})$ are location weighted kernels that introduce spatial dependence in the model (we always use a Gaussian kernel). Lastly, $\mu_h^*(\mathbf{s}_i) \stackrel{iid}{\sim} N(\mu_0, \sigma_0^2)$ for $h = 1, \dots, k_n$ and $\mu_0 \sim N(0, 10^2)$, $\sigma_0 \sim UN(0, 10)$.

The SSB is included because it is algorithmically similar to the sPPM and was fit using the R function provided by Reich and Fuentes (2007). Since the function only admits models that do not include likelihood spatial structure, to make comparisons valid, we do not incorporate spatial structure in (18). The `spBayes` package in R (Finley and Banerjee 2013) was used to fit the SR model.

We considered the following four factors:

1. Number of clusters (either 1 or 4);
2. Distribution of ϵ_i (either $N(0, \sigma^2)$ or $0.5N(0, \sigma^2) + 0.5N(1, \sigma^2)$ with $\sigma^2 = 0.1$ in both cases);
3. Value of M ;
4. Shapes of clusters (square or random).

The first factor was considered to assess clustering accuracy. Note the the sPPM and SSB will, by definition, create spatially referenced clusters, so we do not expect high clustering accuracy when the number of clusters is 1. But including this level will allow us to assess the CPS when the true data generating mechanism is much simpler. Factors 2 and 4 are included to assess robustness of predictions and of goodness-of-fit against possible model perturbations. Finally, factor 3 is included to investigate M 's influence on inference and clustering.

To create synthetic data we employed the following as a data generating mechanism:

$$y(\mathbf{s}_i) = \mu_{c_i}^*(\mathbf{s}_i) + x(\mathbf{s}_i)\beta + \theta(\mathbf{s}_i) + \epsilon(\mathbf{s}_i)$$

$$\boldsymbol{\theta} = [\theta(\mathbf{s}_1), \dots, \theta(\mathbf{s}_n)] \sim GP(\mathbf{0}, \tau^2 \mathbf{H}(\phi)).$$

An exponential covariance function with $\tau^2 = 2$ and $\phi = 6$ was used to create $\mathbf{H}(\phi)$. Locations $(\mathbf{s}_1, \dots, \mathbf{s}_n)$ were generated in two ways. In the first method, we set $\mathbf{s}_i \stackrel{iid}{\sim} UN(0, 1) \times UN(0, 1)$ with clusters being created by partitioning the \mathbb{R}^2 simplex into four equal area squares and assigning \mathbf{s}_i accordingly. For the second method, we set $\mathbf{s}_i \stackrel{iid}{\sim} \sum_{k=1}^4 0.25N(m_k, s_k^2)$. Locations were generated from the mixture using the **MixSim** R function (see Melnykov et al. 2012) which requires specifying the average component overlap (which we set to 5%) and a maximum overlap (which we set to 20%). For data containing four clusters, values of the cluster specific intercepts were $\boldsymbol{\mu}^* = (0, 1, -1, -2)$. We set $\beta = 1$ for all data sets and used $UN(0, 10)$ to generate x values. To obtain point estimates for ρ we employed the least squares procedure proposed in Dahl (2006).

For each combination of factor levels $D = 100$ data sets containing 100 training and 100 testing observations were generated. For each data set, the SSB, SR and sPPM procedures were fit to data by collecting 1000 MCMC iterates after discarding the first 1000 as burn in. Results for $M = 0.01$, $M = 0.1$, and $M = 1.0$ are presented in tabular form and can be found in Tables 3 and 4 (results for other values of M are provided in the Supplementary Materials file (Page and Quintana 2015b)). The table columns correspond to the following:

- RAND represents the adjusted Rand index which measures proximity of estimated partition to the true partition. An adjusted Rand index close to 1 indicates a good match between estimated and true partition. In Table 3, we report the adjusted Rand index averaged over the $D = 100$ data sets.

Error	Cluster	Method	$M = 0.01$			$M = 0.1$			$M = 1.0$		
			RAND	LPML	MSPE	RAND	LPML	MSPE	RAND	LPML	MSPE
Gaussian	Square	CPS $C_{1\alpha=1}$	0.16	-178.07	2.43	0.09	-172.61	2.45	0.05	-169.73	2.75
		CPS $C_{1\alpha=2}$	0.18	-180.28	2.27	0.12	-179.75	2.30	0.07	-182.67	2.51
		CPS C_2	0.43	-187.08	2.30	0.34	-183.87	2.24	0.18	-181.98	2.28
		CPS C_3	0.35	-186.96	2.34	0.41	-184.62	2.26	0.45	-182.30	2.21
		CPS C_4	0.47	-182.97	2.24	0.43	-180.28	2.17	0.30	-178.41	2.18
		SSB	0.03	-201.48	3.36	0.05	-200.71	3.39	0.11	-199.41	3.41
		SR	-	-2804.15	22.02	-	-2501.09	21.93	-	-2669.12	22.27
		CPS $C_{1\alpha=1}$	0.27	-176.76	2.28	0.14	-173.83	2.39	0.07	-166.78	2.55
Irregular	Irregular	CPS $C_{1\alpha=2}$	0.29	-176.51	2.21	0.18	-175.87	2.24	0.09	-179.16	2.39
		CPS C_2	0.49	-185.54	2.38	0.42	-183.50	2.39	0.25	-182.84	2.40
		CPS C_3	0.55	-185.03	2.31	0.59	-183.57	2.26	0.60	-183.95	2.22
		CPS C_4	0.66	-180.88	2.19	0.67	-179.31	2.14	0.56	-179.00	2.11
		SSB	0.13	-195.17	3.11	0.12	-191.86	2.95	0.15	-191.94	2.97
		SR	-	-2632.71	21.36	-	-2267.52	21.62	-	-2460.53	21.04
		CPS $C_{1\alpha=1}$	0.16	-176.90	2.36	0.09	-172.04	2.43	0.05	-169.89	2.62
		CPS $C_{1\alpha=2}$	0.18	-178.30	2.28	0.12	-179.01	2.33	0.06	-183.29	2.61
Mixture	Square	CPS C_2	0.43	-184.89	2.34	0.34	-183.22	2.24	0.18	-182.86	2.37
		CPS C_3	0.36	-184.87	2.34	0.41	-184.05	2.26	0.45	-184.06	2.32
		CPS C_4	0.47	-181.30	2.23	0.42	-179.60	2.17	0.28	-179.54	2.28
		SSB	0.03	-199.44	3.38	0.04	-200.29	3.37	0.09	-198.99	3.49
		SR	-	-2400.44	21.91	-	-2383.69	22.17	-	-2320.54	22.40
		CPS $C_{1\alpha=1}$	0.27	-176.37	2.27	0.17	-176.83	2.46	0.07	-170.99	2.61
		CPS $C_{1\alpha=2}$	0.28	-175.93	2.10	0.17	-176.68	2.21	0.09	-181.69	2.42
		CPS C_2	0.49	-185.90	2.35	0.44	-183.85	2.34	0.26	-185.31	2.44
Irregular	Irregular	CPS C_3	0.56	-186.57	2.21	0.58	-185.33	2.21	0.62	-185.93	2.26
		CPS C_4	0.69	-181.95	2.06	0.63	-179.69	2.08	0.56	-180.84	2.14
		SSB	0.14	-194.10	2.95	0.14	-191.98	2.85	0.15	-193.66	3.01
		SR	-	-2420.39	21.71	-	-2412.06	21.67	-	-2445.62	21.61

Table 3: Simulation study results when data are generated with four clusters.

- LPML is the log-pseudo-marginal likelihood which is a goodness-of-fit metric (see Christensen et al. 2011) that takes into account model complexity. The values in the two tables are the average LPML over the 100 data sets.
- MSPE represents the mean squared prediction error defined as

$$\frac{1}{100} \sum_{i=1}^{100} (Y_p(\mathbf{s}_{di}) - \hat{Y}_p(\mathbf{s}_{di}))^2$$

where i indexes the 100 testing observations ($Y_p(s)$) and

$$\hat{Y}_p(\mathbf{s}_{di}) = E(Y_p(\mathbf{s}_{di}) | \mathbf{Y}(\mathbf{s})).$$

The values found in Tables 3 and 4 are the MSPE averaged over the 100 data sets.

Table 3 provides results for data that contain four clusters. First notice that when $\alpha = 1$ for C_1 the LPML associated with CPS declines as M increases, but prediction accuracy and Rand index values improve. This indicates that M must be small when $\alpha = 1$ for C_1 or CPS tends to overfit by creating many clusters. For C_3 it appears that the opposite is true. Setting $\alpha = 2$ for C_1 seems to reduce overfitting as model fit is slightly worse but out of sample prediction greatly improves. It seems like C_4 is the best at making accurate predictions regardless of the value of M , but selecting an appropriate M is clearly cohesion dependent (something we explore more in the Supplementary Material (Page and Quintana 2015b)). Interestingly CPS (and SSB) predict slightly better when error is a mixture and clusters are not regular. All that said, perhaps the main take home message is that CPS produces more accurate predictions and better data fit relative to SSB and SR for almost all data generating scenarios and cohesions.

Table 4 provides results for data with no clusters. Notice that we do not report the Rand index in this scenario as the CPS and SSB by construction create clusters. Because of this the one cluster partition is not recovered well. That said, this scenario allows us to assess over-fit properties as the data structure is much simpler. It turns out that the model fits associated with data that contain no clusters are similar to those produced with data that contain four clusters. However, the MSPE values are slightly better (which was expected). Generally speaking, it appears that CPS continues to perform well relative to SSB for each of the cohesions and SR.

4.2 Application: Chilean Standardized Testing

Over the past 25 years Chile's Ministry of Education has established a national large-scale standardized test called SIMCE (Sistema de Medición de la Calidad de la Educación, System Measurement of Quality of Education). It was introduced during the later part of the 1980s and since then has continually grown in scope and scale and is now a key component of Chilean educational policies (Meckes and Carrasco, 2010; Manzi and Preiss, 2013). During the early part of the 1980s, education was privatized in Chile allowing parents a great deal of flexibility in selecting a school to send their

Error	Cluster	Method	$M = 0.01$		$M = 0.1$		$M = 1.0$	
			LPML	MSPE	LPML	MSPE	LPML	MSPE
Gaussian	Square	CPS $C_{1,\alpha=1}$	-173.75	2.06	-175.93	2.10	-173.36	2.09
		CPS $C_{1,\alpha=2}$	-173.19	1.98	-174.29	2.01	-173.82	2.02
		CPS C_2	-180.46	2.05	-179.80	2.08	-177.92	2.06
		CPS C_3	-180.75	2.11	-181.88	2.12	-179.72	2.07
		CPS C_4	-177.55	2.05	-177.68	2.03	-175.50	1.99
		SSB	-179.29	2.11	-181.11	2.16	-177.97	2.15
		SR	-2275.31	19.99	-2803.85	19.59	-2504.16	20.11
		CPS $C_{1,\alpha=1}$	-171.41	1.90	-171.85	1.96	-170.50	2.00
Irregular	Irregular	CPS $C_{1,\alpha=2}$	-169.76	1.81	-170.60	1.86	-169.65	1.91
		CPS C_2	-177.02	1.91	-177.87	1.95	-174.27	1.97
		CPS C_3	-180.50	1.94	-179.84	1.98	-177.64	1.97
		CPS C_4	-176.59	1.86	-175.91	1.88	-173.01	1.88
		SSB	-176.59	2.03	-178.35	2.10	-175.69	2.09
		SR	-1913.70	19.58	-1902.62	20.13	-2115.85	19.77
		CPS $C_{1,\alpha=1}$	-176.19	2.08	-173.84	2.11	-174.30	2.04
		CPS $C_{1,\alpha=2}$	-175.03	2.00	-172.48	2.02	-174.22	1.97
Mixture	Square	CPS C_2	-181.95	2.08	-178.33	2.08	-177.63	2.01
		CPS C_3	-182.77	2.12	-179.73	2.12	-179.47	2.03
		CPS C_4	-180.04	2.05	-176.24	2.03	-176.07	1.95
		SSB	-180.15	2.16	-178.39	2.19	-178.56	2.08
		SR	-2470.62	19.47	-2776.41	19.97	-2532.80	19.16
		CPS $C_{1,\alpha=1}$	-172.84	1.98	-170.92	1.87	-170.66	1.92
		CPS $C_{1,\alpha=2}$	-171.30	1.90	-169.55	1.79	-169.85	1.84
		CPS C_2	-179.63	2.00	-175.63	1.87	-174.98	1.90
Irregular	Irregular	CPS C_3	-181.69	2.03	-178.91	1.88	-178.72	1.91
		CPS C_4	-178.14	1.96	-174.74	1.81	-173.71	1.83
		SSB	-178.50	2.10	-176.28	1.97	-176.03	2.04
		SR	-2040.83	19.94	-2291.49	19.47	-1847.70	20.25

Table 4: Simulation study results when data are generated with one cluster.

children. One of the purported roles of SIMCE is to aid parents in making this decision. In addition to administrating the exam, other socio-economic variables are recorded. Among them is mother's education level which is known to influence individual SIMCE scores. Therefore, we include mother's education as a covariate in modeling.

We briefly note that accommodating spatial dependence in education studies has only very recently been considered. In fact, the one article we found is Neelon et al. (2014). They explore regional differences in end-of-grade test scores in North Carolina using county level data. This was done by modeling reading and math scores jointly through a fairly sophisticated joint conditional autoregressive model.

We were given access to individual 2011 SIMCE 4th grade math scores. To simplify the analysis, instead of analyzing individual test scores and mother's education level, we compute school specific averages for both variables. The longitude and latitude of each

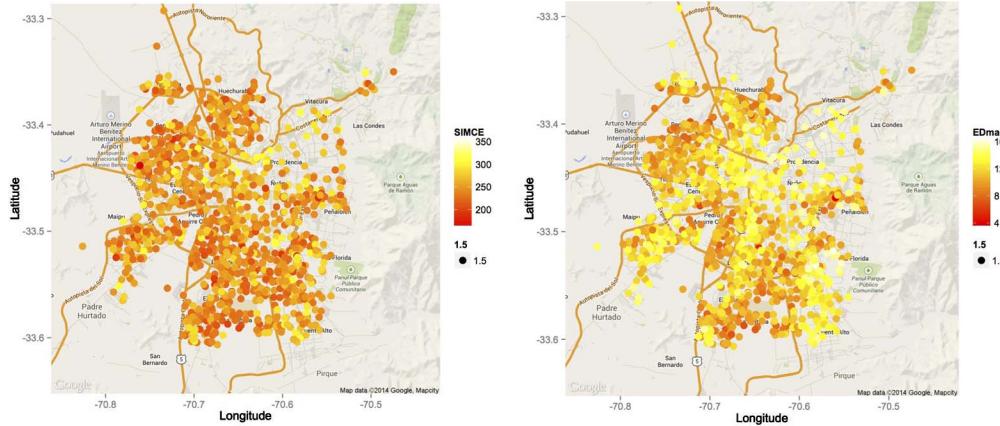


Figure 5: Spatial plots of SIMCE math scores and mother education level. The left figure corresponds with average SIMCE math scores, while the right average mothers education level.

school was recorded and we focus only on the 1215 schools that are located in the greater Santiago area. Figure 5 provides a spatial plot for both SIMCE and mother's education values. Notice that schools in the north east part of the city tend to have higher SIMCE scores than those in the south and west. Mother's education level also varies spatially with lower levels generally appearing in the west and south of Santiago. An exploratory analysis was performed to investigate spatial structures in the SIMCE data, results of which are provided in the Supplementary Material (Page and Quintana 2015b).

To demonstrate the flexibility of pairing the sPPM with a variety of likelihoods, in what follows we detail and compare three reasonable models that could be proposed for the SIMCE data. All the models that will be described were fit to data by collecting 1000 MCMC iterates after discarding the first 10,000 as burn-in and thinning by 20. Convergence was monitored graphically. The MCMC chains mixed reasonably well and converged quickly.

To assess out of sample prediction, we divided the 1215 schools into 600 training observations and 615 testing observations. This partitioning of the data also facilitated a cross validation study (see Supplementary Material (Page and Quintana 2015b)) that in addition to information gleaned from the simulation study resulted in setting M equal to 5×10^{-5} , 0.1, 1.0, and 0.5 for cohesions 1–4, respectively. For C_1 both $\alpha = 1$ and $\alpha = 2$ were considered, but only results from $\alpha = 1$ are reported as $\alpha = 2$ produced very similar fits. The tuning parameters associated with other cohesions are those employed previously.

Conditional Model

In order to compare fits and predictions associated with sPPM to those of SSB, our first modeling approach is to model SIMCE scores conditional on mother education

level with spatial structure in the prior only. This model corresponds to the CPS model of Section 4.1.

In this section and Section 4.3, we consider standardized (zero mean, unit variance) SIMCE and mother’s education scores in addition to the original (unstandardized) scores. A result of standardizing the response is that only residual structure remains in the data and therefore clustering is completely driven by spatial dependence. This is not the case for unstandardized data, where the mean structure also influences clustering. Both are considered to highlight differences between sPPM and SSB.

To compare model fit we once again employ LPML (see Christensen et al. 2011), but now also include $MSE = \frac{1}{n} \sum_{i=1}^n (y(\mathbf{s}_i) - \hat{y}(\mathbf{s}_i))^2$ and the Watanabe–Akaike information criterion (WAIC) which is a fairly new hierarchical model selection metric advocated in Gelman et al. (2014). The MSPE associated with the 615 testing observations is also provided under the “MSPE” column of Table 5.

For the standardized data where only residual spatial structure influences clustering, it appears that CPS fits the data better than SSB. Additionally, CPS appears to make more accurate predictions compared to SSB with C_4 producing the most accurate. The CPS under C_1 clearly fits the data best and produces competitive predictions.

For the unstandardized data, CPS under C_1 continues to fit data better than other procedures, but SSB is more competitive in terms of model fit compared to the standardized data scenario. This results from the fact that sPPM places more mass on spatially connected partitions than SSB. Since in this example mother’s education level explains much of the spatial dependence in SIMCE scores, spatially connected partitions do not provide as much benefit (compare this to the results found in Section 4.3). For this reason SSB produces competitive fits relative to CPS. That said, CPS still does much better compared to SSB in out of sample prediction regardless of cohesion. Thus it seems that spatially connected partitions improve predictions regardless of which structure (mean or covariance) is explaining the spatial dependence.

For the CPS procedure predicting an average SIMCE score for a completely new school requires knowing the new school’s location and mother’s education level. One

	Procedure	WAIC	LPML	MSE	MSPE
Standardized	CPS C_1	2113.64	-1314.21	0.12	0.533
	CPS C_2	2420.56	-1358.97	0.21	0.535
	CPS C_3	2739.73	-1364.31	0.48	0.538
	CPS C_4	2706.71	-1361.58	0.40	0.516
	SSB	2733.40	-1387.91	0.48	0.536
Unstandardized	CPS C_1	4592.05	-2517.47	29.77	366.45
	CPS C_2	4694.46	-2598.12	43.00	399.76
	CPS C_3	4673.26	-2561.56	54.37	369.77
	CPS C_4	4669.92	-2565.81	41.09	361.89
	SSB	4625.12	-2539.50	44.30	574.97

Table 5: Model fit comparisons associated with SIMCE test score data for sPPM and SSB.

approach would be to discretize mother's education into, say, three levels and create a predictive map for each one. An alternative approach would be to first predict mother's education level for the new school, then use the predicted mother's education level as covariate to predict SIMCE. Using the later approach, the 600 training observations, and a regular grid of locations that belonged to the convex hull created by the observed school locations, we predict SIMCE scores by first predicting mother's education level using a model similar to CPS but free of covariates, i.e., $z(\mathbf{s}_i) | \rho, \boldsymbol{\mu}^*, \sigma^2 \sim N(\mu_{c_i}^*(\mathbf{s}_i), \sigma^2)$ where $z(\mathbf{s}_i)$ denotes mother's education level at the i th new school. The predictive map of mother's education values and SIMCE scores is provided in Figure 6 (we only report predictions from C_1 as the others were similar). The predicted values of mother's education level and SIMCE math scores are completely plausible and the resulting spatial structure follows the general socio-economic spatial distribution that is known to exist in Santiago.

Joint Model

Making predictions with the previous model is somewhat awkward as mother's education needs to be either fixed or predicted using a completely different model. A more natural and coherent modeling approach for this application would be to model SIMCE scores and mother's education jointly as both could be thought of as random quantities (in this section we only consider standardized scores). To demonstrate flexibility in which sPPM can be incorporated in modeling and because comparisons to the SSB are not available for the joint model, we include spatial structure in the likelihood which amounts to using a simple coregionalization model (Banerjee et al. 2015, Chapter 9). Now let $\mathbf{y}(\mathbf{s}_i) = [y_1(\mathbf{s}_i), y_2(\mathbf{s}_i)]'$ denote the i th school's average SIMCE score and mother's education level and consider the following data model

$$\mathbf{y}(\mathbf{s}_i) = \boldsymbol{\mu}_{c_i}^*(\mathbf{s}_i) + \boldsymbol{\theta}(\mathbf{s}_i) + \boldsymbol{\epsilon}(\mathbf{s}_i), \quad i = 1, \dots, n, \quad (19)$$

where $\boldsymbol{\mu}_{c_i}^*(\mathbf{s}_i) = [\mu_{1c_i}^*(\mathbf{s}_i), \mu_{2c_i}^*(\mathbf{s}_i)]'$ is a cluster specific two-dimensional intercept vector whose spatial structure is guided through an sPPM prior, $\boldsymbol{\theta}(\mathbf{s}_i) = (\theta_1(\mathbf{s}_i), \theta_2(\mathbf{s}_i))'$ is a two-dimensional intercept whose spatial structure is directly incorporated into the likelihood in a manner that will be described shortly, and $\boldsymbol{\epsilon}(\mathbf{s}_i) \sim N_2(\mathbf{0}, \boldsymbol{\Sigma})$ is an error term. $\boldsymbol{\Sigma}$ contains dependence structure between SIMCE scores and mother's education with variances denoted by σ_1^2 and σ_2^2 and covariance $\sigma_{12} = \eta\sigma_1\sigma_2$. For $h = 1, \dots, k_n$ we assume $\boldsymbol{\mu}_h^*(\mathbf{s}_i) \stackrel{iid}{\sim} N_2(\boldsymbol{\mu}_0, \mathbf{T})$. To address spatial structure for each variable and the dependence that may exist between these two spatial processes, instead of modeling $\theta_1(\mathbf{s}_i)$ and $\theta_2(\mathbf{s}_i)$ directly with a Gaussian process we instead introduce $(\tilde{\theta}_j(\mathbf{s}_1), \tilde{\theta}_j(\mathbf{s}_2), \dots, \tilde{\theta}_j(\mathbf{s}_n)) \sim GP(\mathbf{0}, \mathbf{C}_j)$ independently for $j = 1, 2$ and set

$$\begin{pmatrix} \theta_1(\mathbf{s}_i) \\ \theta_2(\mathbf{s}_i) \end{pmatrix} = \mathbf{A} \begin{pmatrix} \tilde{\theta}_1(\mathbf{s}_i) \\ \tilde{\theta}_2(\mathbf{s}_i) \end{pmatrix} \text{ where } \mathbf{A} = \begin{pmatrix} 1 & \gamma \\ \gamma & 1 \end{pmatrix},$$

for $\gamma \in (-1, 1)$. \mathbf{C}_j of the Gaussian process denotes a valid covariance matrix constructed using an exponential covariance function. Thus, the (ℓ, ℓ') th entry of (\mathbf{C}_j) is $(\mathbf{C}_j)_{\ell, \ell'} = \tau_j^2 \exp\{-\phi_j \|\mathbf{s}_\ell - \mathbf{s}_{\ell'}\|\}$. Prior distributions employed are $\tau_j^2 \sim Gamma(1, 1)$,

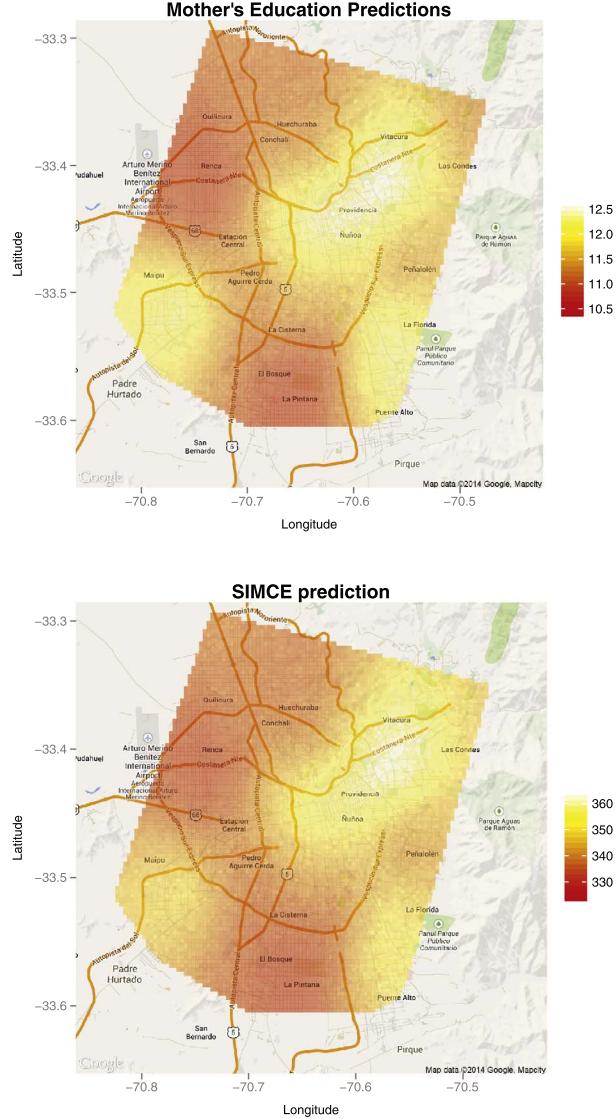


Figure 6: Predictive maps for mother's education and SIMCE scores. The predicted mother's education levels were used to predict SIMCE.

$\phi_j \sim UN(0.5, 30)$ (this implies a $UN(0.1, 6)$ for effective range), $\gamma \sim UN(-1, 1)$, $\mu_0 \sim N_2(\mathbf{0}, 10^2 \mathbf{I})$, $\mathbf{T} \sim IW(2, \mathbf{I})$, and $\Sigma \sim IW(2, \mathbf{I})$. We use $IW(\nu, \Lambda)$ to denote an inverse Wishart distribution with scale and matrix parameters ν and Λ .

Under this model prediction of the SIMCE math score for a new school located at s_0 is easily made via $y_1(s_0)|y_2(s_0)$ which has the following form

$$y_1(\mathbf{s}_0) | y_2(\mathbf{s}_0) \sim N\left(\beta_{0c_0}^*(\mathbf{s}_0) + \beta_1^* y_2(\mathbf{s}_0), \sigma_1^2(1 - \eta^2)\right),$$

with $\beta_1^* = \eta \frac{\sigma_1}{\sigma_2}$ and $\beta_{0c_0}^*(\mathbf{s}_i) = \mu_{1c_0}^* + \theta_1(\mathbf{s}_0) - \beta_1^* [\mu_{2c_0}^* + \theta_2(\mathbf{s}_0)]$.

For this procedure to be useful, predictions of $\mu_{1c_0}^*$, $\mu_{2c_0}^*$, $\theta_1(\mathbf{s}_0)$, $\theta_2(\mathbf{s}_0)$, and $y_2(\mathbf{s}_0)$ are needed. Values for μ_1^* and μ_2^* are readily available once c_0 is classified by way of the predictive distribution found in Section 2 of the Supplementary Material (S.1) (Page and Quintana 2015b). Values for $[\theta_1(\mathbf{s}_0), \theta_2(\mathbf{s}_0)]$ are obtained by first predicting $[\tilde{\theta}_1(\mathbf{s}_0), \tilde{\theta}_2(\mathbf{s}_0)]$ from $\tilde{\theta}_1(\mathbf{s}_0) | \tilde{\theta}_1(\mathbf{s}_1), \dots, \tilde{\theta}_1(\mathbf{s}_n)$ and $\tilde{\theta}_2(\mathbf{s}_0) | \tilde{\theta}_2(\mathbf{s}_1), \dots, \tilde{\theta}_2(\mathbf{s}_n)$ independently and then setting $[\theta_1(\mathbf{s}_0), \theta_2(\mathbf{s}_0)]' = \mathbf{A}[\tilde{\theta}_1(\mathbf{s}_0), \tilde{\theta}_2(\mathbf{s}_0)]'$. Finally, using the fact that $y_2(\mathbf{s}_0) \sim N(\mu_{2c_0}^* + \theta_2(\mathbf{s}_0), \sigma_2^2)$ a prediction for $y_2(\mathbf{s}_0)$ is easily obtained. We will refer to the procedure just described as the Joint model with Likelihood Spatial Structure (JLS) model.

JLS can become computationally expensive as the number of schools grows. Incorporating spatial information solely in the prior would radically reduce computation time, but potentially at the cost of model fit. To investigate this trade off, we also consider

$$\begin{aligned} \mathbf{y}(\mathbf{s}_i) | \boldsymbol{\mu}^*, c_i &\stackrel{ind}{\sim} N_2(\boldsymbol{\mu}_{c_i}^*(\mathbf{s}_i), \boldsymbol{\Sigma}) \text{ for } i = 1, \dots, n \text{ and } \boldsymbol{\Sigma} \sim IW(2, \mathbf{I}), \\ \mu_h^* | \boldsymbol{\mu}_0, \mathbf{T} &\stackrel{iid}{\sim} N_2(\boldsymbol{\mu}_0, \mathbf{T}) \text{ with } \mathbf{T} \sim IW(2, \mathbf{I}), \\ \boldsymbol{\mu}_0 &\sim N_2(\mathbf{0}, 10^2 \mathbf{I}), \\ \{c_i\}_{i=1}^n &\sim sPPM. \end{aligned}$$

As in the JLS, predictions at location \mathbf{s}_0 are also easily made via $E[y_1(\mathbf{s}_0) | y_2(\mathbf{s}_0)] = \mu_{1c_0}^*(\mathbf{s}_0) + \eta \frac{\sigma_1}{\sigma_2} [y_2(\mathbf{s}_0) - \mu_{2c_0}^*(\mathbf{s}_0)]$. Values for $\mu_{1c_0}^*(\mathbf{s}_0)$, $\mu_{2c_0}^*(\mathbf{s}_0)$, and $y_2(\mathbf{s}_0)$ are gathered using the procedure described for JLS. We will refer to this model as the Joint model with Prior Spatial Structure (JPS).

Using the same M values as in Section 4.2 we fit JLS and JPS to the training data and carried out prediction using the same grid of points and the testing data. Comparisons of the two joint models regarding model fit and computation time are provided in Table 6. The column ‘‘Clusters’’ is the expected number of clusters *a posteriori* and ‘‘Time’’ is the amount of computing time required to fit models (measured in seconds). MSPE is associated with the 600 testing observations. As expected fits using JLS are much better

Procedure	WAIC	LPML	MSE	MSPE	Clusters	Time
JPS C_1	2312.503	-1383.301	0.380	0.586	35.767	2154
JPS C_2	2569.589	-1438.750	0.415	0.590	34.746	4621
JPS C_3	2778.803	-1447.872	0.482	0.591	8.921	598
JPS C_4	2552.333	-1399.899	0.433	0.600	26.750	1090
JLS C_1	2047.319	-1291.011	0.244	0.574	34.992	38017
JLS C_2	2266.945	-1342.172	0.258	0.569	34.249	41022
JLS C_3	2553.984	-1376.176	0.365	0.573	6.789	38538
JLS C_4	2273.479	-1331.949	0.334	0.606	26.952	37565

Table 6: Model fit comparisons for the JPS and JLS models fit to the SIMCE education data set.

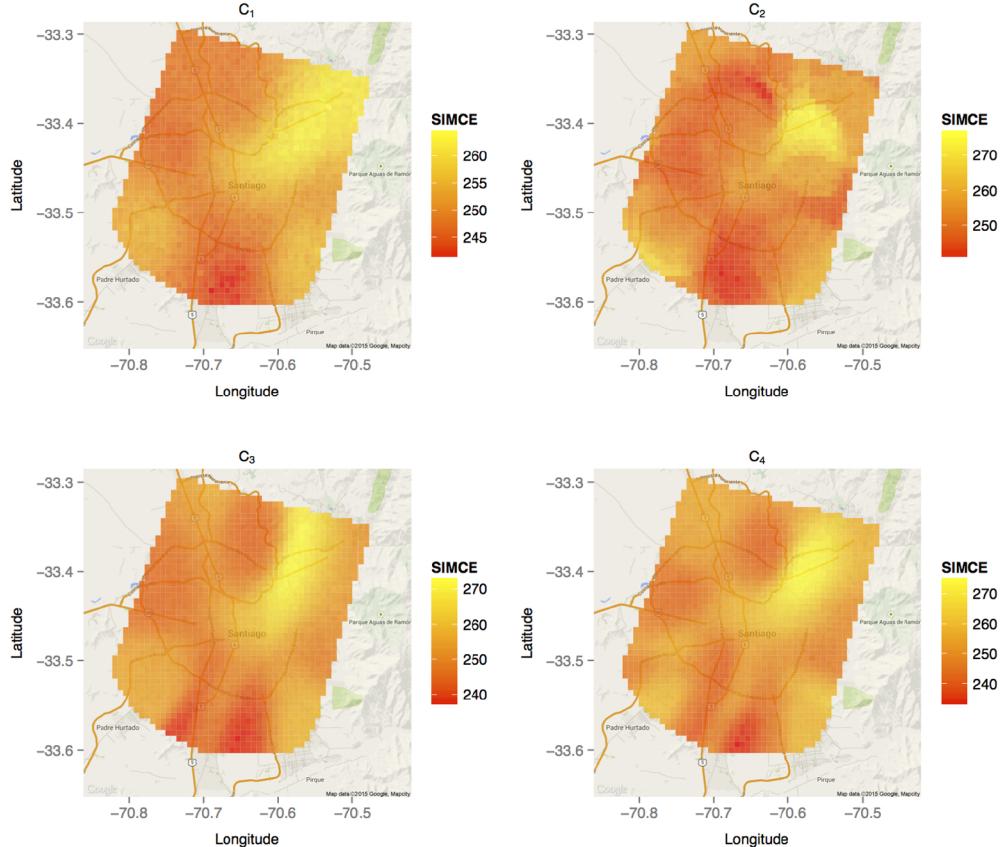


Figure 7: Predictive maps associated with JPS for each of the four cohesion functions.

for all cohesion functions but at a substantial computational cost. However, JPS out of sample predictions are fairly competitive to those from JLS and may be considered if a timely answer is needed.

Maps associated with predictions made using JPS and JLS are provided in Figures 7 and 8. For JPS the four cohesions produce fairly different predictive surfaces, while for JLS the surfaces are very similar among the four cohesions. This illustrates that including spatial structure in the likelihood greatly impacts the predictive maps. For both procedures, the predictive maps identify the same general areas that contain higher SIMCE scores, but changes in SIMCE scores as a function of space are far more pronounced for JLS. This may be indicating that predictions are more local for JLS relative to JPS.

4.3 Scallops Data Application

The Chilean Education example was a very natural application to show the utility of the sPPM. Here we provide a more standard application (the scallops dataset found

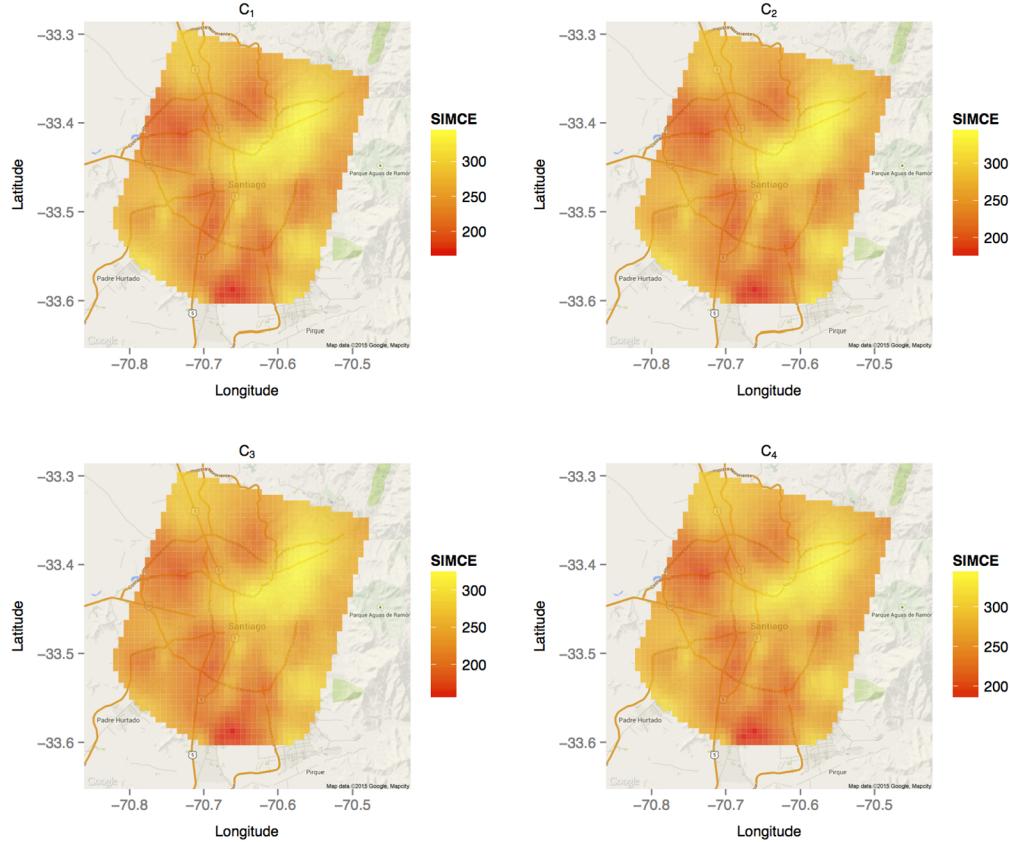


Figure 8: Predictive maps associated with JLS for each of the four cohesion functions.

in Banerjee et al. 2015) that further illustrates sPPM's utility (but only considering C_1 and C_4). In this data set, the total scallop catch was measured at 148 locations in the New York/New Jersey Bight. Figure 9 displays the raw data with the circle circumference proportional to total catch amount. Letting $z(s_i)$ be the total scallop catch at location i , we follow suggestion of Banerjee et al. (2015) and model the log-transformed total scallop catch $y(s_i) = \log(z(s_i) + 1)$.

Since there are no covariates, we employ the same model as that used to model mother's education level of the Chilean Education example of Section 4.2. Specifically, after introducing cluster labels c_1, \dots, c_n , the hierarchical model we employ is

$$\begin{aligned}
 y(s_i) | c_i, \mu_{c_i}^*(s_i), \sigma^2 &\stackrel{\text{ind}}{\sim} N(\mu_{c_i}^*(s_i), \sigma^2) \text{ and } \sigma \sim UN(0, 10), \\
 \mu_j^*(s_i) &\stackrel{iid}{\sim} N(\mu_0, \sigma_0^2) \text{ for } j = 1, \dots, k_n \text{ and} \\
 \mu_0 &\sim N(0, 10^2), \sigma_0 \sim UN(0, 10), \\
 \{c_i\}_{i=1}^n &\sim sPPM.
 \end{aligned} \tag{20}$$



Figure 9: Locations at which scallop catch amounts of New York/New Jersey coast were recorded. Total catch is proportional to circle circumference.

	Procedure	LPML	MSE	WAIC
Standardized	sPPM C_1	-140.01	0.06	182.31
	sPPM C_4	-138.66	0.09	196.98
	SSB	-167.37	0.11	243.22
Unstandardized	sPPM C_1	-271.51	0.38	424.64
	sPPM C_4	-250.96	0.49	430.11
	SSB	-304.74	1.47	584.96

Table 7: Model fit comparisons associated with SIMCE test score data for sPPM and SSB.

We compare model fits to that of the SSB process for standardized and original log transformed catch totals. Following suggestions found in Section 3 of the Supplementary Material (Page and Quintana 2015b), we set $M = 1$ for C_4 and $M = 0.0001$ for C_1 . The tuning parameters associated with C_4 are those used in the main article and $\alpha = 1$ for C_1 . For the SSB $M = 1$.

We use the same model fit metrics as before (LPML, MSE, and WAIC) to compare model fits and Table 7 contains the results. sPPM outperforms SSB for both cohesions much more dramatically for this application than that of the Chilean education data. This holds true regardless of whether data have been standardized or not (data not shown). Further, it appears that the sPPM with C_1 fits the data best. (Predictive maps associated with this application are provided in the Supplementary Material (Page and Quintana 2015b).)

5 Conclusions

We have proposed a general procedure that extends PPMs to a spatial setting. Thus providing a mechanism to directly model the partitioning of locations into spatially

dependent clusters. As a result, incorporating fairly sophisticated spatial structures in modeling can be done straightforwardly. The cohesion function of the sPPM affords a great deal of flexibility regarding the type of spatial clusters that may be considered. We provided four possibilities, two that were motivated from a spatial perspective (C_1, C_2) and two that were motivated by the appealing properties of the covariate dependent PPM (C_3, C_4). In the absence of information associated with tuning parameters, C_4 tends to provide better out of sample predictions and therefore is always a sensible choice. However, it does favor partitions with more clusters relative to C_3 . On the other hand, C_1 performs very well if adequately tuned (e.g., selecting values for M and α via cross-validation). If areal data are available, then C_2 would be the more natural choice. Investigating the influence that tuning parameters have on predictions is a topic of ongoing research. The four cohesions that we have proposed are certainly not exhaustive. Other cohesion functions can be developed that produce different types of spatial structures.

The simulation study and application showed that the methodology is particularly well suited for predictions. Being able to incorporate spatial information in the prior and likelihood allows for added flexibility in how spatial structure is modeled, providing the added benefit of capturing local structure. Exactly how to join local spatial structure so that global maps are smooth and continuous (if so desired) is also a topic of ongoing research. In the case that a measured response is not Gaussian (or even continuous), it should be straightforward to link the sPPM with a non-Gaussian likelihood and employ alternative approaches, such as, e.g., those based on generalized linear models. Finally, although not explicitly considered, including covariate information in the clustering mechanism, in addition to spatial information, should be a natural extension of work developed in Müller et al. (2011).

Appendix A: Marginal Correlation Proof

We provide a detailed proof of Propositions 3.2 and 3.3. The proof of Proposition 3.1 follows very similar arguments.

A.1 Proof of Proposition 3.2

Proof. From the law of total covariance

$$\begin{aligned}
\text{cov}(y_i, y_j) &= \text{cov}_{\rho, \beta, \theta}[E(y_i|\rho, \beta, \theta), E(y_j|\rho, \beta, \theta)] + E_{\rho, \beta, \theta}[\text{cov}(y_i, y_j, |\rho, \beta, \theta)] \\
&= E_{\rho, \beta, \theta}[(\mathbf{x}'_i \boldsymbol{\beta}_{c_i}^* + \theta_i)(\mathbf{x}'_j \boldsymbol{\beta}_{c_j}^* + \theta_j)] - E_{\rho, \beta, \theta}[\mathbf{x}'_i \boldsymbol{\beta}_{c_i}^* + \theta_i]E_{\rho, \beta, \theta}[\mathbf{x}'_j \boldsymbol{\beta}_{c_j}^* + \theta_j] + 0 \\
&= E_{\rho, \beta, \theta}[(\mathbf{x}'_i \boldsymbol{\beta}_{c_i}^*)(\mathbf{x}'_j \boldsymbol{\beta}_{c_j}^*) + (\mathbf{x}'_i \boldsymbol{\beta}_{c_i}^*)\theta_j + \theta_i(\mathbf{x}'_j \boldsymbol{\beta}_{c_j}^*) + \theta_i\theta_j] \\
&\quad - E_{\rho, \beta}[\mathbf{x}'_i \boldsymbol{\beta}_{c_i}^*]E_{\rho, \beta}[\mathbf{x}'_j \boldsymbol{\beta}_{c_j}^*] \\
&= E_{\rho, \beta}[(\mathbf{x}'_i \boldsymbol{\beta}_{c_i}^*)(\mathbf{x}'_j \boldsymbol{\beta}_{c_j}^*)] + E_{\theta}[\theta_i\theta_j] - E_{\rho, \beta}[\mathbf{x}'_i \boldsymbol{\beta}_{c_i}^*]E_{\rho, \beta}[\mathbf{x}'_j \boldsymbol{\beta}_{c_j}^*]
\end{aligned}$$

$$\begin{aligned}
&= \sum_{\rho} E_{\beta} [tr\{\boldsymbol{\beta}_{c_i}^* \mathbf{x}_i \mathbf{x}_j' \boldsymbol{\beta}_{c_j}^*\}] Pr(\rho) \\
&\quad - \left(\sum_{\rho} E_{\beta} [\mathbf{x}_i' \boldsymbol{\beta}_{c_i}^*] Pr(\rho) \right) \left(\sum_{\rho} E_{\beta} [\mathbf{x}_j' \boldsymbol{\beta}_{c_j}^*] Pr(\rho) \right) + cov(\theta_i, \theta_j) \\
&= \sum_{\rho} E_{\beta} [tr\{\mathbf{x}_i \mathbf{x}_j' \boldsymbol{\beta}_{c_j}^* \boldsymbol{\beta}_{c_i}^{*'}\}] Pr(\rho) \\
&\quad - \left(\sum_{\rho} \mathbf{x}_i' \boldsymbol{\mu} Pr(\rho) \right) \left(\sum_{\rho} \mathbf{x}_j' \boldsymbol{\mu} Pr(\rho) \right) + cov(\theta_i, \theta_j) \\
&= \sum_{\rho: c_i = c_j} tr\{\mathbf{x}_i \mathbf{x}_j' (\mathbf{T} + \boldsymbol{\mu} \boldsymbol{\mu}')\} Pr(\rho) \\
&\quad + \sum_{\rho: c_i \neq c_j} tr\{\mathbf{x}_i \mathbf{x}_j' (\boldsymbol{\mu} \boldsymbol{\mu}')\} Pr(\rho) - \boldsymbol{\mu}' \mathbf{x}_i \mathbf{x}_j' \boldsymbol{\mu} + cov(\theta_i, \theta_j) \\
&= \mathbf{x}_j' \mathbf{T} \mathbf{x}_i \sum_{\rho: c_i = c_j} Pr(\rho) + cov(\theta_i, \theta_j) \\
&= \mathbf{x}_j' \mathbf{T} \mathbf{x}_i Pr(c_i = c_j) + \lambda^2 (H(\phi))_{i,j}.
\end{aligned}$$

Now using the law of total variance,

$$\begin{aligned}
var(y_i) &= E_{\rho, \beta, \theta} [var(y_i | \rho, \beta, \theta)] + var_{\rho, \beta, \theta} [E(y_i | \rho, \beta, \theta)] \\
&= E_{\rho, \beta, \theta} [\sigma^2] + var_{\rho, \beta, \theta} [\mathbf{x}_i' \boldsymbol{\beta}_{c_i}^* + \theta_i] \\
&= \sigma^2 + \lambda^2 + \mathbf{x}_i' \mathbf{T} \mathbf{x}_i.
\end{aligned}$$

Using $corr(y_i, y_j) = \frac{cov(y_i, y_j)}{\sqrt{var(y_i)} \sqrt{var(y_j)}}$ completes the proof. \square

A.2 Proof of Proposition 3.3

Proof. Following similar arguments from the previous proof,

$$\begin{aligned}
cov(y_i, y_j) &= cov_{\rho, \beta, \theta} [E(y_i | \rho, \beta, \theta), E(y_j | \rho, \beta, \theta)] + E_{\rho, \beta, \theta} [cov(y_i, y_j, | \rho, \beta, \theta)] \\
&= \mathbf{x}_i' \mathbf{T} \mathbf{x}_j + \sum_{\rho: c_i = c_j} cov(\theta_i, \theta_j) Pr(\rho) + \sum_{\rho: c_i \neq c_j} cov(\theta_i, \theta_j) Pr(\rho) \\
&= \mathbf{x}_i' \mathbf{T} \mathbf{x}_j + \sum_{\rho: c_i = c_j} cov(\theta_i, \theta_j) Pr(\rho) \\
&= \mathbf{x}_i' \mathbf{T} \mathbf{x}_j + \sum_{h=1}^{k_n} \sum_{\rho: c_i = c_j = h} \lambda_h^2 (H(\phi_h))_{i,j} Pr(\rho) \\
&= \mathbf{x}_i' \mathbf{T} \mathbf{x}_j + \sum_{h=1}^{k_n} \lambda_h^2 (H(\phi_h))_{i,j} \sum_{\rho: c_i = c_j = h} Pr(\rho) \\
&= \mathbf{x}_i' \mathbf{T} \mathbf{x}_j + \sum_{h=1}^{k_n} \lambda_h^2 (H(\phi_h))_{i,j} Pr(c_i = c_j = h).
\end{aligned}$$

And now using the law of total variance,

$$\begin{aligned}
 var(y_i) &= E_{\rho, \beta, \theta}[var(y_i|\rho, \beta, \theta)] + var_{\rho, \beta, \theta}[E(y_i|\rho, \beta, \theta)] \\
 &= \sigma^2 + \mathbf{x}'_i \mathbf{T} \mathbf{x}_i + \sum_{\rho} var_{\theta}(\theta_i) Pr(\rho) \\
 &= \sigma^2 + \mathbf{x}'_i \mathbf{T} \mathbf{x}_i + \sum_{h=1}^{k_n} var(\theta_i) \sum_{\rho: c_i = h} Pr(\rho) \\
 &= \sigma^2 + \mathbf{x}'_i \mathbf{T} \mathbf{x}_i + \sum_{h=1}^{k_n} \lambda_h^{2*} Pr(c_i = h).
 \end{aligned}$$

Using $corr(y_i, y_j) = \frac{cov(y_i, y_j)}{\sqrt{var(y_i)}\sqrt{var(y_j)}}$ completes the proof. \square

Supplementary Material

Supplementary Materials for the article titled “Spatial Product Partition Models” (DOI: [10.1214/15-BA971SUPP](https://doi.org/10.1214/15-BA971SUPP); .pdf).

References

- Banerjee, S., Carlin, B. P., and Gelfand, A. E. (2015). *Hierarchical Modeling and Analysis for Spatial Data*. Boca Raton, Florida: Chapman & Hall/CRC, 2nd edition. [MR3362184](#). 266, 287, 291
- Barry, D. and Hartigan, J. A. (1992). “Product Partition Models for Change Point Problems.” *The Annals of Statistics*, 20: 260–279. [MR1150343](#). doi: <http://dx.doi.org/10.1214/aos/1176348521>. 268
- Blei, D. M. and Frazier, P. I. (2011). “Distant dependent Chinese restaurant processes.” *Journal of Machine Learning Research*, 12: 2461–2488. [MR2834504](#). 267
- Christensen, R., Johnson, W., Branscum, A., and Hanson, T. E. (2011). *Bayesian Ideas and Data Analysis*. Texts in Statistical Science Series. CRC Press, Boca Raton, FL. An introduction for scientists and statisticians. [MR2682928](#). 283, 286
- Dahl, D. B. (2006). “Model-Based Clustering for Expression Data via a Dirichlet Process Mixture Model.” In: Vannucci, M., Do, K. A., and Müller, P. (eds.), *Bayesian Inference for Gene Expression and Proteomics*, 201–218. Cambridge University Press. 281
- Denison, D. G. T. and Holmes, C. C. (2001). “Bayesian Partitioning for Estimating Disease Risk.” *Biometrics*, 57: 143–149. [MR1833300](#). doi: <http://dx.doi.org/10.1111/j.0006-341X.2001.00143.x>. 267, 270
- Diggle, P. J. (2014). *Statistical Analysis of Spatial and Spatio-temporal Point Patterns*, volume 128 of *Monographs on Statistics and Applied Probability*. CRC Press, Boca Raton, FL, third edition. [MR3113855](#). 267

- Duan, J. A., Guindani, M., and Gelfand, A. E. (2007). “Generalized Spatial Dirichlet Process Models.” *Biometrika*, 94: 809–825. [MR2416794](#). doi: <http://dx.doi.org/10.1093/biomet/asm071>. 267
- Dunson, D. B. and Park, J.-H. (2008). “Kernel Stick-Breaking Processes.” *Biometrika*, 95: 307–323. [MR2521586](#). doi: <http://dx.doi.org/10.1093/biomet/asn012>. 267
- Finley, A. O. and Banerjee, S. (2013). *spBayes: Univariate and Multivariate Spatial-temporal Modeling*. R package version 0.3-8. <http://CRAN.R-project.org/package=spBayes>. 280
- Gelfand, A., Diggle, P., Guttorp, P., and Fuentes, M. (2010). *Handbook of Spatial Statistics*. Chapman & Hall/CRC Handbooks of Modern Statistical Methods. Taylor & Francis. <http://books.google.es/books?id=Xf4les1PDzsC>. [MR2761512](#). doi: <http://dx.doi.org/10.1201/9781420072884>. 266
- Gelfand, A. E., Kottas, A., and MacEachern, S. N. (2005). “Bayesian Nonparametric Spatial Modeling With Dirichlet Process Mixing.” *Journal of the American Statistical Association*, 100: 1021–1035. [MR2201028](#). doi: <http://dx.doi.org/10.1198/016214504000002078>. 267
- Gelman, A., Hwang, J., and Vehtari, A. (2014). “Understanding predictive information criteria for Bayesian models.” *Statistics and Computing*, 24(6): 997–1016. [MR3253850](#). doi: <http://dx.doi.org/10.1007/s11222-013-9416-2>. 286
- Ghosh, S., Ungureanu, A. B., Sudderth, E. B., and Blei, D. (2011). “Spatial distance dependent Chinese restaurant processes for image segmentation.” In: Shawe-Taylor, J., Zemel, R., Bartlett, P., Pereira, F., and Weinberger, K. (eds.), *Advances in Neural Information Processing Systems 24*, 1476–1484. <http://papers.nips.cc/paper/4589-repulsive-mixtures.pdf>. 267
- Griffin, J. E. and Steel, M. F. J. (2006). “Order-Based Dependent Dirichlet Processes.” *Journal of the American Statistical Association*, 101: 179–194. [MR2268037](#). doi: <http://dx.doi.org/10.1198/016214505000000727>. 267
- Hartigan, J. A. (1990). “Partition Models.” *Communications in Statistics, Part A – Theory and Methods*, 19: 2745–2756. [MR1088047](#). doi: <http://dx.doi.org/10.1080/03610929008830345>. 266, 268
- Hegarty, A. and Barry, D. (2008). “Bayesian Disease Mapping Using Product Partition Models.” *Statistics in Medicine*, 27: 3868–3893. [MR2526613](#). doi: <http://dx.doi.org/10.1002/sim.3253>. 268
- Kang, J., Zhang, N., and Shi, R. (2014). “A Bayesian nonparametric model for spatially distributed multivariate binary data with application to a multidrug-resistant tuberculosis (MDR-TB) study.” *Biometrics*, 70(4): 981–992. [MR3295759](#). doi: <http://dx.doi.org/10.1111/biom.12198>. 267, 271
- Knorr-Held, L. and Raßer, G. (2000). “Bayesian Detection of Clusters and Discontinuities in Disease Maps.” *Biometrics*, 56: 13–21. 268
- Lawson, A. B. (2013). *Bayesian Disease Mapping: Hierarchical Modeling in Spatial Epidemiology*. Chapman and Hall/ CRC, second edition. [MR2484272](#). 269

- Lawson, A. B. and Denison, D. G. T. (2002). *Spatial Cluster Modeling*. Chapman and Hall/ CRC. [MR2015027](#). doi: <http://dx.doi.org/10.1201/9781420035414>. 267
- Lee, D., Rushworth, A., and Sahu, S. K. (2014). “A Bayesian Localized Conditional Autoregressive Model for Estimating the Health Effects of Air Pollution.” *Biometrics*, 70: 419–429. [MR3258046](#). doi: <http://dx.doi.org/10.1111/biom.12156>. 267
- Li, P., Banerjee, S., Hanson, T. A., and McBean, A. M. (2015). “Bayesian Hierarchical Models for Detecting Boundaries in Areally Referenced Spatial Datasets.” *Statistica Sinica*, 25: 385–402. [MR3328821](#). 267
- Manzi, J. and Preiss, D. (2013). “Educational Assessment and Educational Achievement in South America.” In: Hattie, J. and Anderman, E. M. (eds.), *International Guide to Student Achievement*, Chapter 9. Taylor and Friends. 283
- Meckes, L. and Carrasco, R. (2010). “Two decades of SIMCE: An overview of the National Assessment System in Chile.” *Assessment in Education: Principles, Policy and Practice*, 17: 233–248. 283
- Melnykov, V., Chen, W. C., and Maitra, R. (2012). “MixSim: An R Package for Simulating Data to Study Performance of Clustering Algorithms.” *Journal of Statistical Software*, 51(12): 1–25. <http://www.jstatsoft.org/v51/i12/>. 281
- Monteiro, J. V. D., Assunção, R. M., and Loschi, R. H. (2011). “Product partition models with correlated parameters.” *Bayesian Analysis*, 6: 691–726. [MR2869962](#). 268
- Müller, P., Quintana, F., and Rosner, G. L. (2011). “A Product Partition Model With Regression on Covariates.” *Journal of Computational and Graphical Statistics*, 20(1): 260–277. [MR2816548](#). doi: <http://dx.doi.org/10.1198/jcgs.2011.09066>. 268, 271, 293
- Neal, R. M. (2000). “Markov Chain Sampling Methods for Dirichlet Process Mixture Models.” *Journal of Computational and Graphical Statistics*, 9: 249–265. [MR1823804](#). doi: <http://dx.doi.org/10.2307/1390653>. 280
- Neelon, B., Gelfand, A. E., and Miranda, M. L. (2014). “A Multivariate Spatial Mixture Model for Areal Data: Examining Regional Differences in Standardized Test Scores.” *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 63(5): 737–761. [MR3269410](#). doi: <http://dx.doi.org/10.1111/rssc.12061>. 284
- Page, G. L. and Quintana, F. A. (2015a). “Predictions Based on the Clustering of Heterogeneous Functions via Shape and Subject-Specific Covariates.” *Bayesian Analysis*, 10: 379–410. 268
- Page, G. L. and Quintana, F. A. (2015b). “Supplementary materials for the article titled ‘‘Spatial Product Partition Models’’.” *Bayesian Analysis* (2015). doi: <http://dx.doi.org/10.1214/15-BA971SUPP>. 268, 272, 280, 281, 283, 285, 289, 292
- Papageorgiou, G., Richardson, S., and Best, N. (2014). “Bayesian non-parametric models for spatially indexed data of mixed type.” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*. doi: <http://dx.doi.org/10.1111/rssb.12097>. 267

- Park, J.-H. and Dunson, D. B. (2010). “Bayesian Generalized Product Partition Model.” *Statistica Sinica*, 20: 1203–1226. [MR2730180](#). 268, 270, 271
- Petrone, S., Guindani, M., and Gelfand, A. E. (2009). “Hybrid Dirichlet Mixture Models for Functional Data.” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 94: 755–782. [MR2750094](#). doi: <http://dx.doi.org/10.1111/j.1467-9868.2009.00708.x>. 267
- Quintana, F. A., Müller, P., and Papoila, A. L. (2015). “Cluster-Specific Variable Selection for Product Partition Models.” *Scandinavian Journal of Statistics*. doi: <http://dx.doi.org/10.1111/sjos.12151>. 272
- Reich, B. J. and Bondell, H. D. (2011). “A Spatial Dirichlet Process Mixture Model for Clustering Population Genetics Data.” *Biometrics*, 67: 381–390. [MR2829007](#). doi: <http://dx.doi.org/10.1111/j.1541-0420.2010.01484.x>. 267
- Reich, B. J. and Fuentes, M. (2007). “A Multivariate Semiparametric Bayesian Spatial Modeling Framework for Hurricane Surface Wind Fields.” *The Annals of Applied Statistics*, 1: 249–264. [MR2393850](#). doi: <http://dx.doi.org/10.1214/07-AOAS108>. 267, 280
- Ren, L., Du, L., Carin, L., and Dunson, D. B. (2011). “Logistic Stick-Breaking Processes.” *Journal of Machine Learning Research*, 12: 203–239. [MR2773552](#). 267
- Robert, C. P. and Casella, G. (2010). *Introducing Monte Carlo methods with R*. Use R! Springer, New York. [MR2572239](#). doi: <http://dx.doi.org/10.1007/978-1-4419-1576-4>. 273
- Schabenberger, O. and Gotway, C. A. (2005). *Statistical Methods for Spatial Data Analysis*. Chapman & Hall/CRC. [MR2134116](#). 266
- Sethuraman, J. (1994). “A constructive definition of Dirichlet priors.” *Statistica Sinica*, 4: 639–650. [MR1309433](#). 267
- Wall, M. M. (2004). “A Close Look at the Spatial Structure Implied by the CAR and SAR Models.” *Journal of Statistical Planning and Inference*, 121: 311–324. [MR2038824](#). doi: [http://dx.doi.org/10.1016/S0378-3758\(03\)00111-3](http://dx.doi.org/10.1016/S0378-3758(03)00111-3). 265

Acknowledgments

The first author was partially funded by grant FONDECYT 11121131, and the second author was partially funded by grant FONDECYT 1141057. The authors thank Carolina Flores for granting access to the Chilean education data whose collection was partially funded by the ANILLO Project SOC 1107 Statistics for Public Policy in Education from the Chilean Government.